

4. Instruction Set Architecture – from C to assembly – Basic blocks

EECS 370 – Introduction to Computer Organization – Winter 2015

Robert Dick, Andrew Lukefahr, and Satish Narayanasamy

EECS Department
University of Michigan in Ann Arbor, USA

© Dick-Lukefahr-Narayanasamy, 2015

The material in this presentation cannot be
copied in any form without our written permission

Announcements

- ❑ HW-1 due on 1/22

Recap

❑ LC-2K

- Instructions in LC2K
- Assembling
- Addressing: only base+displacement


❑ ARM

- Arithmetic, logic instructions
- Many loads and stores in ARM
- Addressing: base+displacement and base+register

❑ Slides from the previous lecture

Class Problem 2

- Show the C and ARM assembly for extracting the value in bits 15:10 from a 32-bit integer variable



Want these bits

Remember each hex digit is 4 bits

ARM Memory Instructions

- ❑ Supports base + displacement and base + register modes only
 - Base is a register
 - Offset is a 12-bit immediate, either positive or negative

- ❑ Format: 2 registers (dest, base) and 12-bit immediate (offset)
 - Example:
ldr r3, [r4, #1000] // load word
Retrieves 32-bit value from memory location (r4 + 1000) and puts the result into r3

Load Instruction Sizes

How much data is retrieved from memory at the given address?

- ❑ `ldr r3, [r4, #1000]`
 - retrieve a word (32 bits) from address (r4+1000)

- ❑ `ldrh r3, [r4, #1000]`
 - retrieve a halfword (16 bits) from address (r4+1000)

- ❑ `ldrb r3, [r4, #1000]`
 - retrieve a byte (8 bits) from address (r4+1000)

Sign/Zero Extension

- ❑ Registers in ARM are 32 bits!
- ❑ So what happens when you load 8 or 16 bits?

- Sign extend if the load is signed



- Zero extend if the load is unsigned



ARM Memory Instructions - Summary

❑ Load instructions

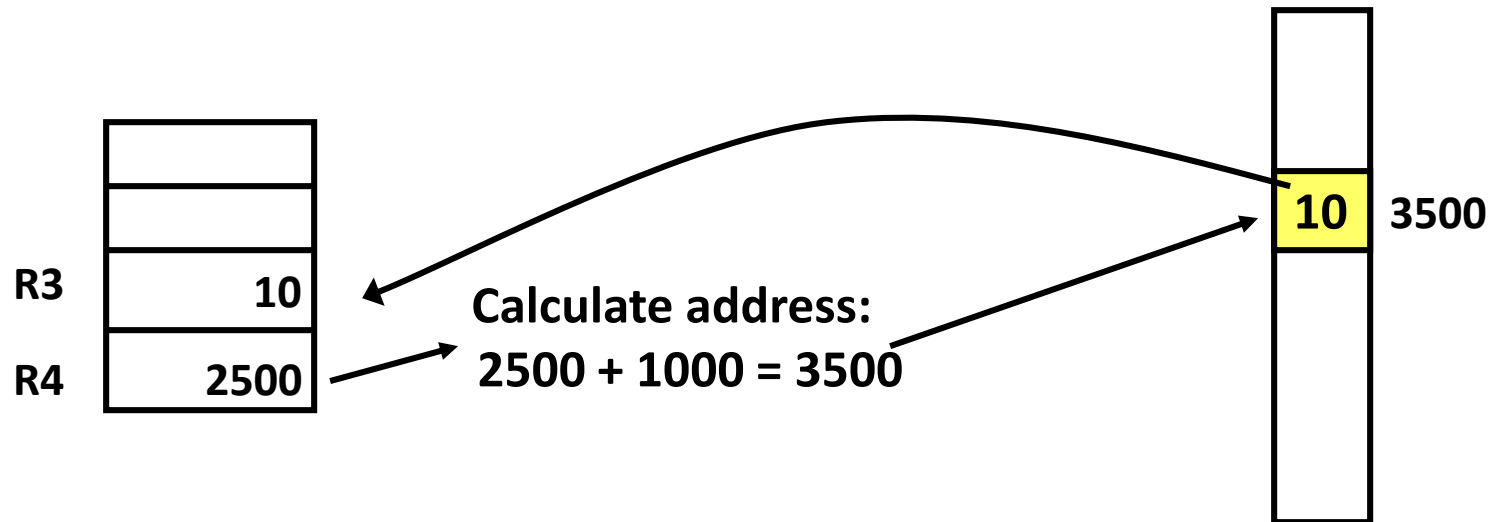
- `ldrsb` `\\` load byte signed (load 8 bits, sign extend)
- `ldrb` `\\` load byte unsigned (load 8 bits, zero extend)
- `ldrsh` `\\` load halfword (load 16 bits, sign extend)
- `ldrh` `\\` load halfword unsigned (load 16 bits, zero extend)
- `ldr` `\\` load word (load 32 bits, no extension)

❑ Store instructions (No sign/zero extension for stores)

- `strb r3, [r4, #1000]` `\\` store 8 LSBs of r3 to M[r4+1000]
- `strh r3, [r4, #1000]` `\\` store 16 LSBs of r3 to M[r4+1000]
- `str r3, [r4, #1000]` `\\` store all 32 bits of r3 to M[r4+1000]


Load Instruction in Action

- ❑ `ldrsb r3, [r4, #1000] // load signed byte`
Retrieves 8-bit value from memory location ($r4+1000$) and puts the result into r3 (sign extended)



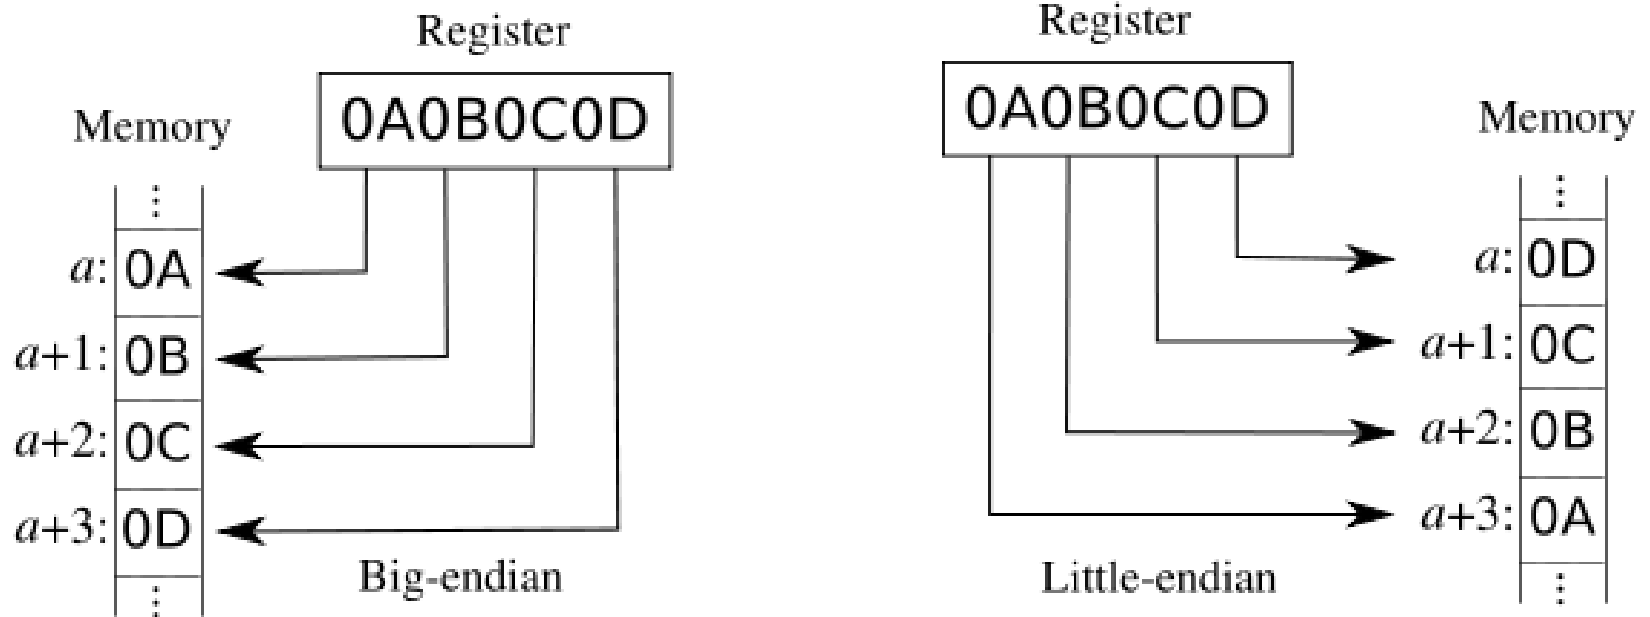
Load Instruction in Action – other example

- ❑ `ldrsb r3, [r4, #1000] // load signed byte`
Retrieves 8-bit value from memory location ($r4+1000$) and puts the result into r3 (sign extended)



Big Endian vs. Little Endian

- Endian-ness: ordering of bytes within a word
 - Little - increasing numeric significance with increasing memory addresses
 - Big – The opposite, most significant byte first
 - The Internet is big endian, x86 and ARM are little endian



Example Code Sequence


What is the final state of memory once you execute the following instruction sequence? (assume r0 has the value of 0)

```
ldr    r4, [r0, #100]
ldrshb r3, [r0, #102]
str    r3, [r0, #100]
strb   r4, [r0, #102]
```

register file



Memory
(each location is 1 byte)




Example Code Sequence – insn 1


What is the final state of memory once you execute the following instruction sequence?

```
ldr      r4, [r0, #100]
ldrshb   r3, [r0, #102]
str       r3, [r0, #100]
strb      r4, [r0, #102]
```

register file



Memory
(each location is 1 byte)



Example Code Sequence – insn 2

What is the final state of memory once you execute the following instruction sequence?

```
ldr    r4, [r0, #100]
ldrsb  r3, [r0, #102]
str     r3, [r0, #100]
strb    r4, [r0, #102]
```




Memory
(each location is 1 byte)

0x02	100
0x03	101
0xFF	102
0x05	103


Example Code Sequence – insn 3

What is the final state of memory once you execute the following instruction sequence?

```
ldr    r4, [r0, #100]
ldrshb r3, [r0, #102]
str    r3, [r0, #100]
strb   r4, [r0, #102]
```




Memory
(each location is 1 byte)




Example Code Sequence – insn 4

What is the final state of memory once you execute the following instruction sequence?

```
ldr    r4, [r0, #100]
ldrshb r3, [r0, #102]
str     r3, [r0, #100]
strb   r4, [r0, #102]
```



Memory
(each location is 1 byte)




Class Problem 3

What is the final state of memory once you execute the following instruction sequence? (assume r0 has the value of 0)

```
ldrh    r3, [r0, #100]
ldrb    r4, [r0, #102]
str      r3, [r0, #100]
strh     r4, [r0, #102]
```




Memory
(each location is 1 byte)




Class Problem 3

What is the final state of memory once you execute the following instruction sequence? (assume r0 has the value of 0)

```
ldrh    r3, [r0, #100]
ldrb     r4, [r0, #102]
str      r3, [r0, #100]
strh     r4, [r0, #102]
```



Memory
(each location is 1 byte)



❑ Slides from the previous lecture ends

Instruction Set Architecture (ISA) Design Lectures

- ❑ Lecture 2: Storage types and addressing modes
- ❑ Lecture 3 : LC-2K and ARM architecture
- ❑ **Lecture 4 : Converting C to assembly – basic blocks**
- ❑ Lecture 5 : Converting C to assembly – functions
- ❑ Lecture 6 : Translation software; libraries, memory layout

Converting C to Assembly

- ❑ Memory layout → memory addresses
- ❑ Branches
- ❑ Procedure calls
- ❑ Expression trees
- ❑ Register allocation

Converting C to assembly – example 1

Write ARM assembly code for the following C expression:

C: `a = b + names[i];`

Assume that **a** is in r1, **b** is in r2, **i** is in r3, and the array **names** starts at address 1000 and holds 32-bit integers

```
mov    r6, #4    // set up integer size
mul    r5, r3, r6 // calculate array offset i*4
ldr    r4, [r5, #1000] // load names[i]
add    r1, r2, r4    // calculate b + names[i]
```

Converting C to assembly – example 2

Write ARM assembly code for the following C expression:

```
class { int a; unsigned char b, c; } y;
```

```
y.a = y.b + y.c;
```

Assume that a pointer to **y** is in r1.

```
ldrb    r2, [r1, #4] // load y.b
```

```
ldrb    r3, [r1, #5] // load y.c
```

```
add     r4, r2, r3 // calculate y.b+y.c
```

```
str     r4, [r1, #0] // store y.a
```

How do you determine the offsets for the class sub-fields?

Calculating Load/Store Addresses for Variables

DATA LAYOUT

```
short a[100];  
char b;  
int c;  
double d;  
short e;  
struct {  
    char f;  
    int g[1];  
    char h;  
} i;
```

Problem: Assume data memory starts at address 100, calculate the total amount of memory needed

$a = 2 \text{ bytes} * 100 = 200$

$b = 1 \text{ byte}$

$c = 4 \text{ bytes}$

$d = 8 \text{ bytes}$

$e = 2 \text{ bytes}$

$i = 1 + 4 + 1 = 6 \text{ bytes}$

total = 221, right or wrong?

Memory layout of variables

- ❑ For ARM, cannot always arbitrarily pack variables into memory
 - Need to worry about alignment
 - An N-byte variable must start at an address A, such that $(A \% N) == 0$
 - Newer ARM processors will perform unaligned loads/stores, but they are VERY SLOW

- ❑ “Golden” rule – Address of a variable is aligned based on the size of the variable
 - **char** is byte aligned (any addr is fine)
 - **short** is half-word aligned (LSBit of addr must be 0)
 - **int** is word aligned (2 LSBit's of addr must be 0)

Structure/Class alignment

- ❑ Each field is laid out in the order it is declared using the Golden Rule for alignment

- ❑ Identify largest field
 - Starting address of overall struct is aligned based on the largest field
 - Size of overall struct is a multiple of the largest field
 - Reason for this is so we can have an array of structs

Structure Example

```
struct {  
    char w;  
    int x[3]  
    char y;  
    short z;  
}
```

The largest field is **int** (4 bytes), hence:



struct size is multiple of 4



struct's starting addr is word aligned

Assume struct starts at location 1000,

char w → 1000

x[0] → 1004-1007, x[1] → 1008 – 1011, x[2] → 1012 – 1015

char y → 1016

short z → 1018 – 1019

Total size = 20 bytes!

Earlier Example – 2nd Try

Assume data memory starts at address 100

```
short a[100];  
char b;  
int c;  
double d;  
short e;  
struct {  
    char f;  
    int g[1];  
    char h;  
} i;
```

→ 200 bytes → 100-299

→ 1 byte → 300-300

→ 4 bytes → 304-307

→ 8 bytes → 312-319

→ 2 bytes → 320-321

→ largest field is 4 bytes → start at 324

→ 1 byte → 324-324

→ 4 bytes → 328 - 331

→ 1 byte → 332-332

→ struct size is 12 bytes, spanning 324 – 335

236 bytes total!!

Class Problem 1

- How much memory is required for the following data, assuming that the data starts at address 200?

```
int a;  
struct {double b, char c, int d} e;  
char *f;  
short g[20];
```

ARM Sequencing Instructions

- ❑ Sequencing instructions change the flow of instructions that are executed
 - This is achieved by modifying the program counter (r15)
- ❑ Conditional branches
 - If (condition_test) goto target_address
 - *condition_test* examines flags from the processor status word (PSR)
 - *target_address* is a 24-bit word displacement on current PC+8

- **cmp r1, r2**
beq label
- if (r1 == r2) then PC = label else PC = PC + 4

ARM Condition Codes Determine Direction of Branch

- ❑ Most arithmetic/logic instructions set condition codes in PSR
 - add, sub, cmp, and, eor, etc...
- ❑ Four primary condition codes evaluated:
 - N – set if the result is **negative** (i.e., bit 31 is non-zero)
 - Z – set if the result is **zero** (i.e., all 32 bits are zero)
 - C – set if last addition/subtraction had a **carry**/borrow out of bit 31
 - V – set if the last addition/subtraction produced an **overflow** (e.g., two negative numbers added together produce a positive result)
- ❑ Branch conditions:

• eq – (Z == 1)	gt – (Z == 0 N == V)
• ne – (Z == 0)	le – (Z == 1 N != V)
• ge – (N == V)	al – 1 (can use shorthand “b label”)
• lt – (N != V)	

Setting the Branch Displacement Field

- ❑ Target address is a 24-bit signed aligned displacement on current PC+8
 - Target = PC + 8 + 4*24_bit_signed_displacement
 - beq 1 // branch **3** instructions ahead if flag Z == 1
 - beq -3 // branch **1** instruction back if flag Z == 1
 - beq -2 // Infinite loop if flag Z == 1

Example code sequence



add
sub
mul
beq
ldr
str
mov

Other Branching Instructions

- `cmp r2, r3` // branch to $4 \times \text{offset} + \text{PC} + 8$ if $r2 \neq r3$
`bne offset`
- `cmp r2, r3` // branch to $4 \times \text{offset} + \text{PC} + 8$ if $r2 < r3$
`blt offset`
- `cmp r2, r3` // branch to $4 \times \text{offset} + \text{PC} + 8$ if $r2 \geq r3$
`bge offset`
- `b offset` // jump to $4 \times \text{offset} + \text{PC} + 8$
// unconditional jump, is taken regardless of PSR flags
- `mov r15, r3` // jump to address in r3 -- when is this useful?
- `bl offset` // put $\text{PC} + 4$ into register r14 (LR) and jump to $4 \times \text{offset} + \text{PC} + 8$

Branch - Example

Convert the following C code into ARM assembly (assume x is in r1, y in r2):

```
int x, y;  
if (x == y)  
    x++;  
(L1) else  
    y++;  
(L2) ...
```

Using Labels

```
cmp r1, r2  
bne L1  
add r1, r1, #1  
b L2  
L1: add r2, r2, #1  
L2: ...
```

Without Labels

```
cmp r1, r2  
bne 1  
add r1, r1, #1  
b 0  
add r2, r2, #1
```

Assemblers must deal with labels and assign displacements

Loop – Example

```
// assume all variables are integers
// i is in r1, start of a is at address 500, sum is in r2
for (i=0 ; i < 100 ; i++) {
    if (a[i] > 0) {
        sum += a[i];
    }
}
```

of branch instructions
= $3 \times 100 = 300$

a.k.a. while-do template

```
        mov     r1, #0
        mov     r4, #400
Loop1:   cmp     r1, r4
        bge     endLoop
        ldr     r5, [r1, #500]
        cmp     r5, #0
        ble     endIf
        add     r2, r2, r5
endIf:   add     r1, r1, #4
        b       Loop1
endLoop:
```

Same Loop, Different Assembly

```
// assume all variables are integers
// i is in r1, start of a is at address 500, sum is in r2
for (i=0 ; i < 100 ; i++) {
    if (a[i] > 0) {
        sum += a[i];
    }
}
```

of branch instructions
= $2 \times 100 + 1 = 201$

a.k.a. do-while template

```
        mov     r1, #0
        mov     r4, #400
        cmp     r1, r4
        bge     endLoop
Loop1:   ldr     r5, [r1, #500]
        cmp     r5, #0
        ble     endlf
        add     r2, r2, r5
endlf:   add     r1, r1, #4
        cmp     r1, r4
        blt     Loop1
endLoop:
```

Class Problem 2

Write the ARM assembly code to implement the following C code:

```
// assume ptr is in r1
// struct {int val; struct node *next;} node;
// struct node *ptr;

if ((ptr != NULL) && (ptr->val > 0))
    ptr->val++;
```