

Basic Processor Design – Introduction to Pipelining

EECS 370 – Introduction to Computer Organization – Winter 2015

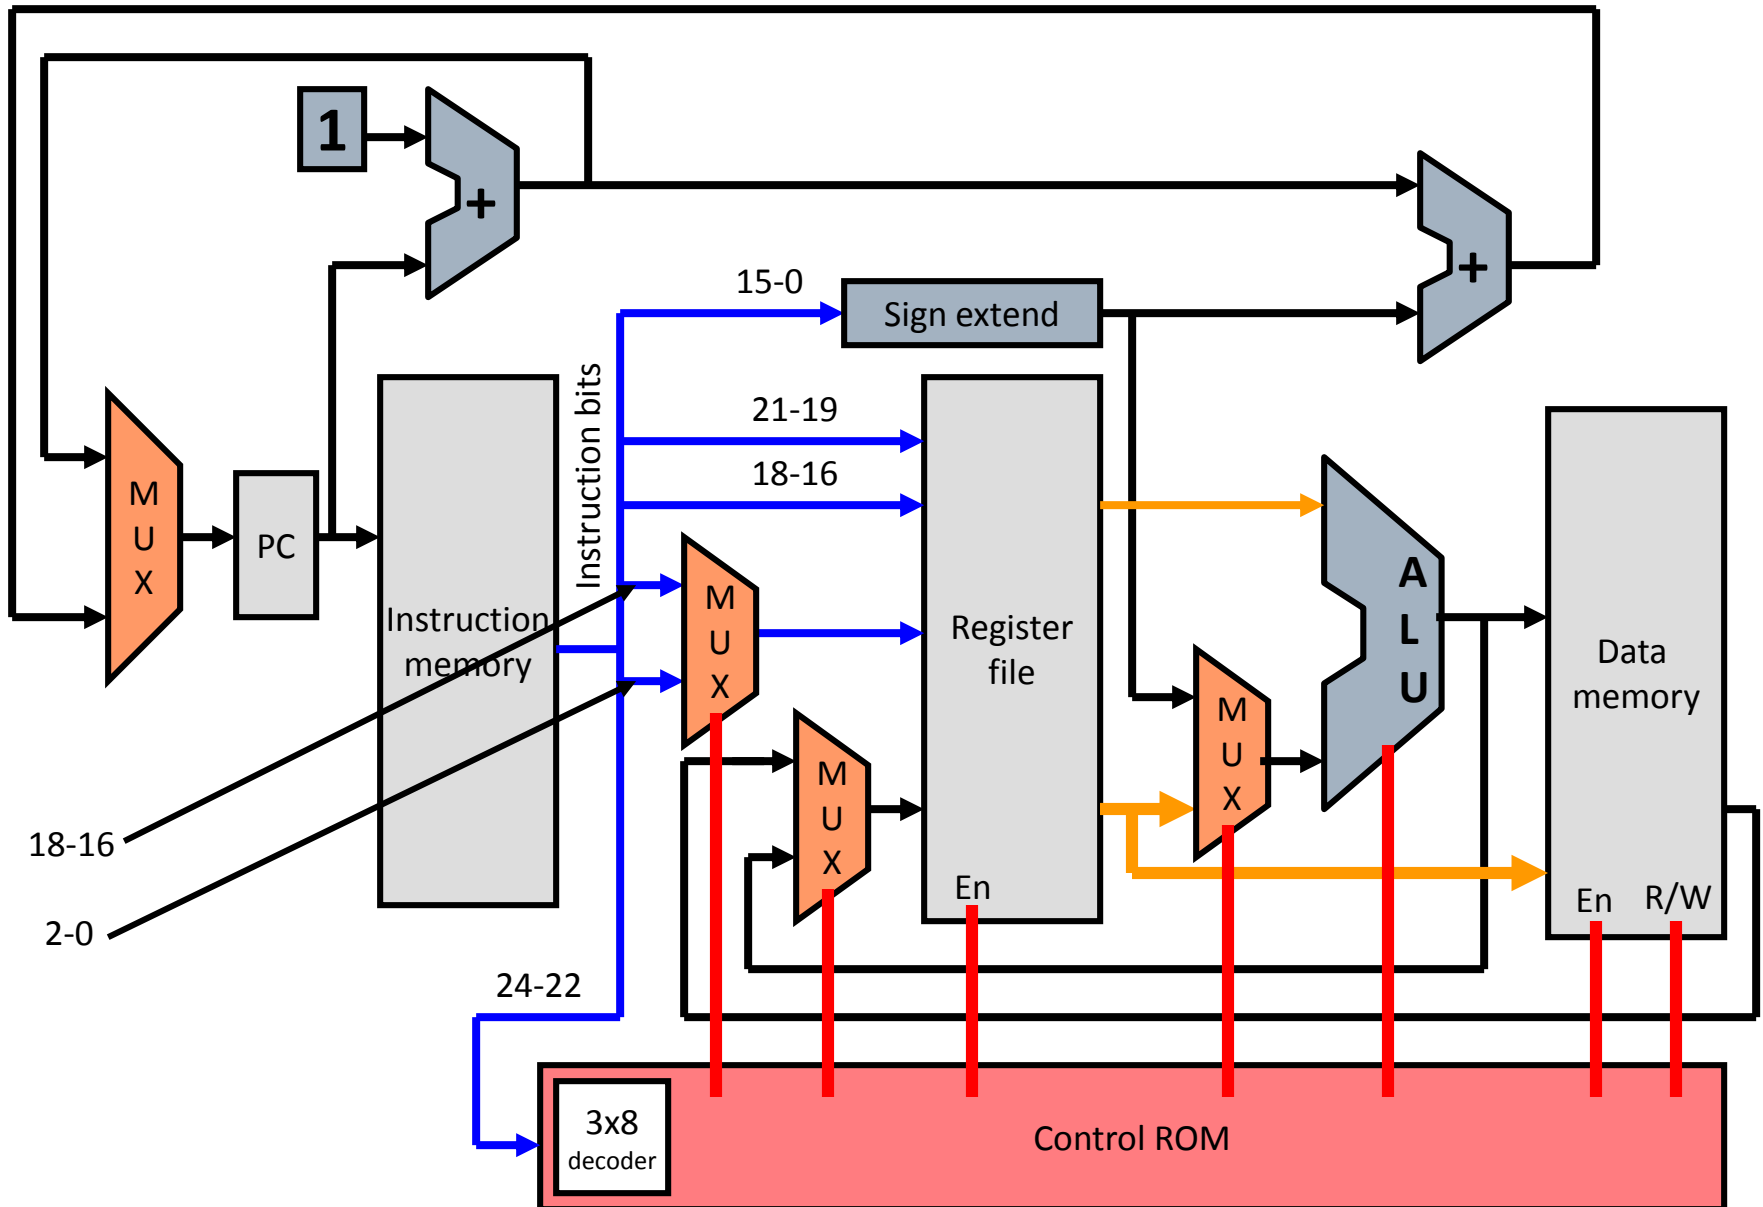
Robert Dick, Andrew Lukefahr, and Satish Narayanasamy

EECS Department
University of Michigan in Ann Arbor, USA

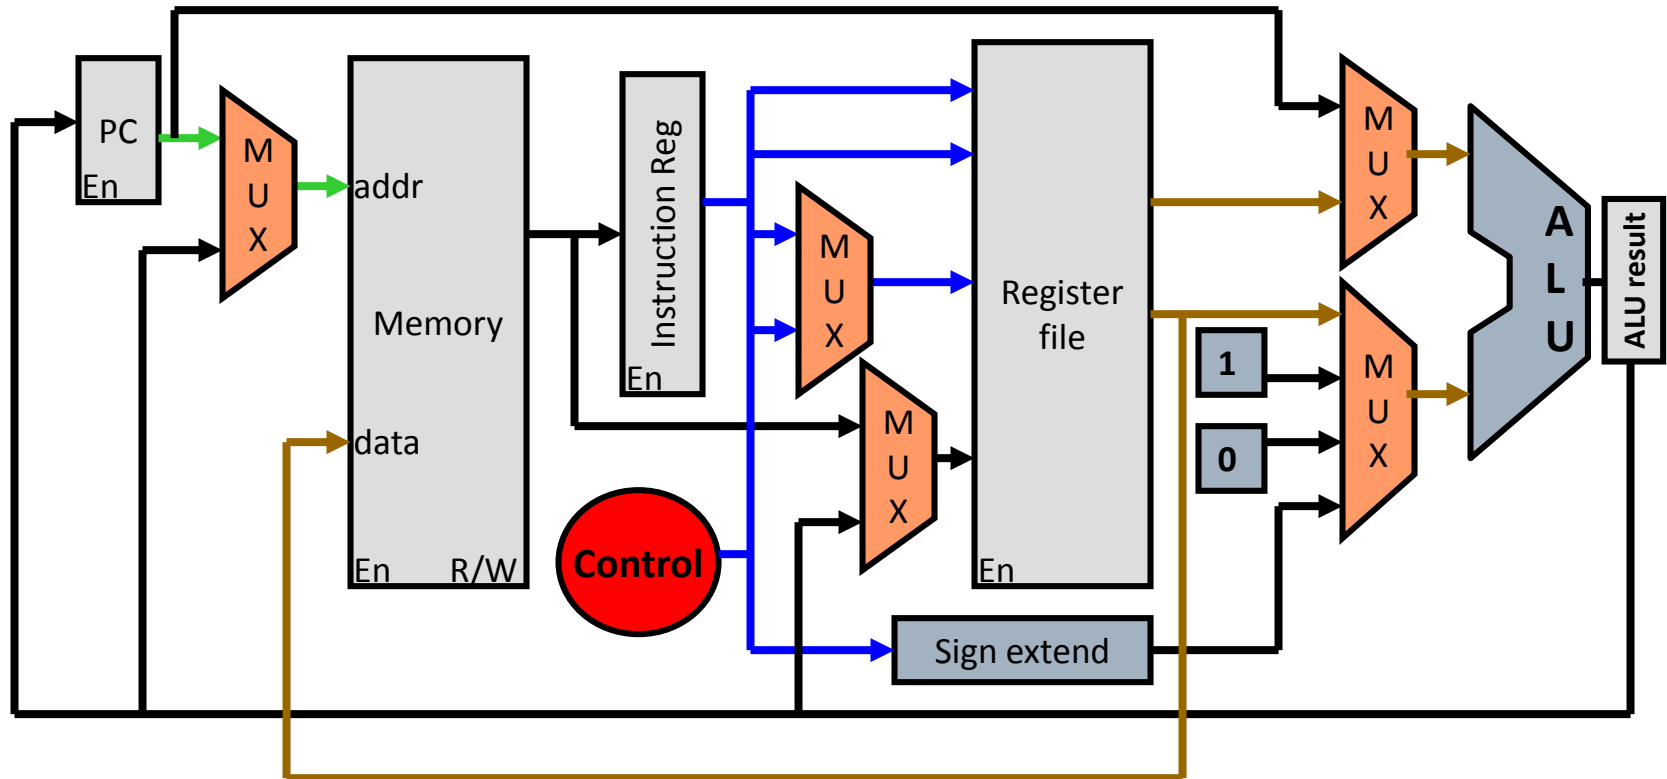
© Dick-Lukefahr-Narayanasamy, 2015

The material in this presentation cannot be
copied in any form without our written permission

Review: Single-Cycle LC2Kx Datapath



Review: Multi-Cycle LC2Kx Datapath



Review: Single and Multicycle Performance

1 ns – Register read/write time

2 ns – ALU/adder

2 ns – memory access

0 ns – MUX, PC access, sign extend, ROM

1. Assuming the above delays, what is the best cycle time that the LC2k multicycle datapath could achieve?

MC: $\text{MAX}(2, 1, 2, 2, 1) = 2\text{ns}$

SC: $2 + 1 + 2 + 2 + 1 = 8\text{ ns}$

2. Assuming the above delays, for a program consisting of 25 LW, 10 SW, 45 ADD, and 20 BEQ, which is faster?

SC: $100\text{ cycles} * 8\text{ ns} = 800\text{ ns}$

MC: $(25*5 + 10*4 + 45*4 + 20*4)\text{cycles} * 2\text{ns} = 850\text{ ns}$


Performance Metrics

1. **Response time:** when is my job done (time)?
 - When will my books arrive from amazon.com?
 - How long will this program/instruction take?

2. **Throughput:** how much work can get done within a specified time (work/time)?
 - How many books will amazon.com sell this week?
 - How many programs/instructions complete per hour?
 - Improved relatively easily by using multiprocessors.

Relating response time to throughput

More throughput \leftrightarrow lower response time (*Little's law*)



Performance Metrics – Execution Time

- ❑ Response time for a program is its **execution time**

Execution time (for an application):

= total instructions executed x CPI x clock period

- Called the “Iron Law” of performance

- ❑ CPI = **avg** number of clock cycles per instruction *for an application*
- ❑ For multi-cycle processor implementations we need:
 - Cycles necessary for each type of instruction
 - Mix of instructions executed in the application (dynamic instruction execution profile)

How far have we come?

❑ Single-cycle processor implementations

- $\text{CPI} = ?$
- $\text{clock period} = ?$

❑ Multi-cycle processor implementations

- $\text{CPI} = ?$
- $\text{clock period} = ?$

❑ Next step: improve CPI without impacting clock period

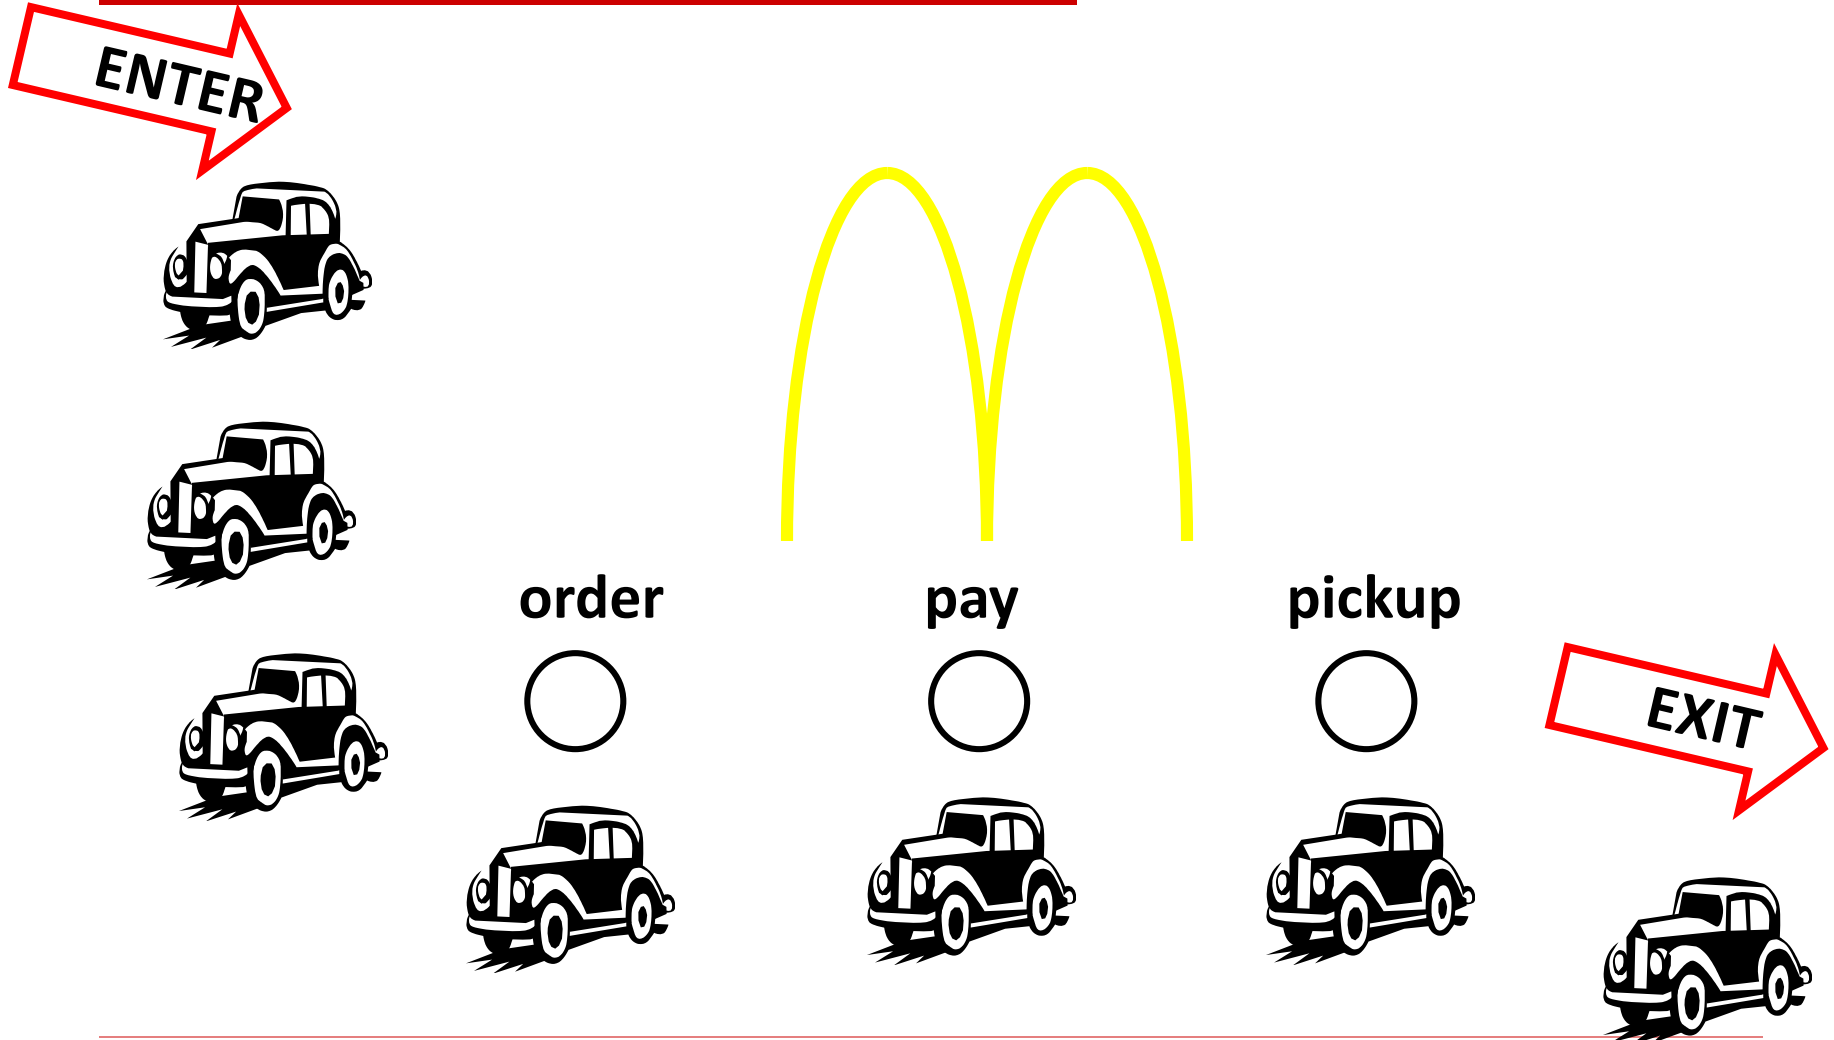
- The easiest thing to do is to work on multiple instructions at the same time.

Pipelining

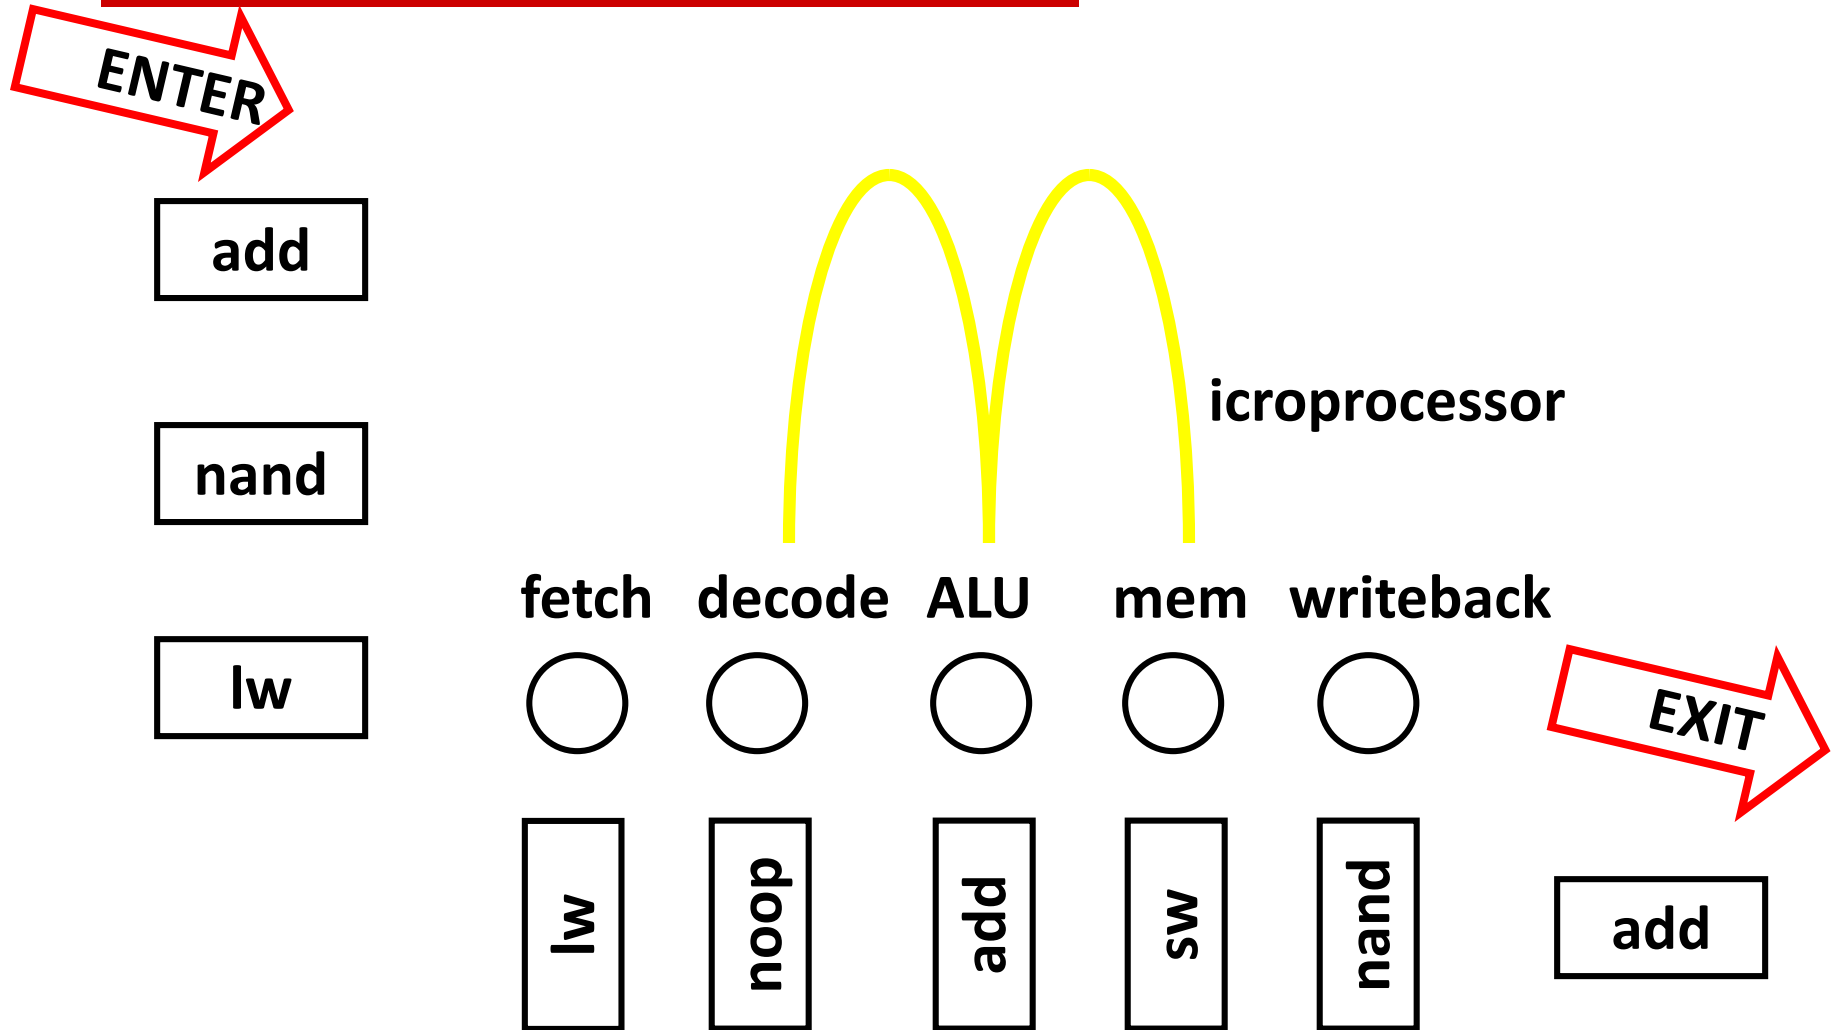
❑ Want to execute an instruction?

- Build a processor (multi-cycle)
- Find instructions
- Line up instructions (1, 2, 3, ...)
- Overlap execution
 - Cycle #1: Fetch 1
 - Cycle #2: Decode 1 Fetch 2
 - Cycle #3: ALU 1 Decode 2 Fetch 3
 -
- This is called pipelining instruction execution.
- Used extensively for the first time on IBM 360 (1960s).
- CPI approaches 1.

Pipelining



Pipelining



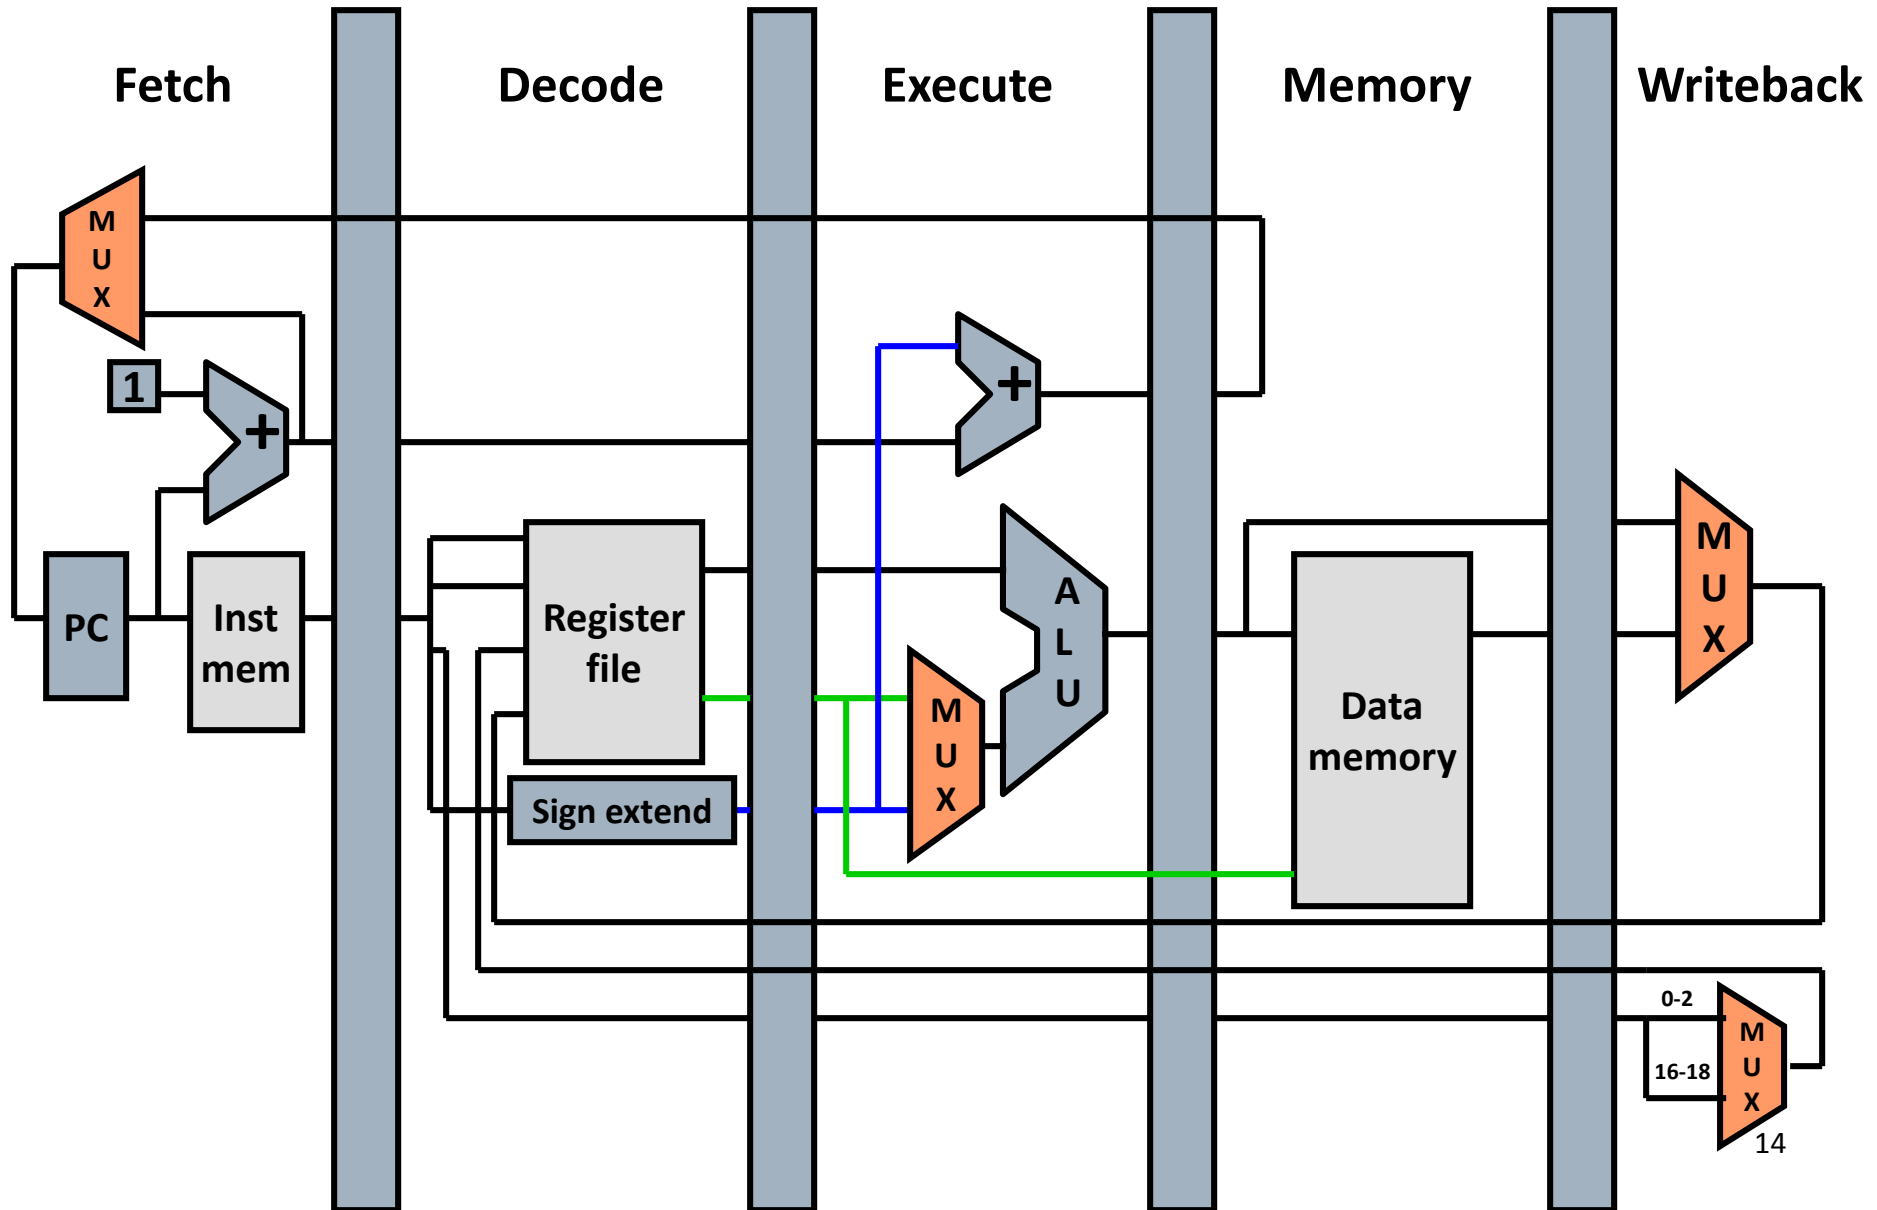
Pipelining Today

- ❑ Execute as many instructions at the same time as possible.
 - Pipelining: 12-20+ cycles
 - Multiple pipelines
- ❑ Pentium:
 - 2 pipelines, 5 cycles each (10 instructions “in flight”)
- ❑ Pentium Pro/II/III
 - 3 pipelines (kinda), 12 cycles each (kinda)
 - Instructions can execute out of their original program order
- ❑ Pentium IV
 - 4 pipelines, 20 cycles deep
 - Prescott: 4 pipelines, 31 cycles deep (could be clocked up to 8 GHz with special cooling)
- ❑ Core i7 (Nehalem)
 - 4 pipelines, 16 cycles deep

Pipelined implementation of LC2Kx

- ❑ Break the execution of the instruction into cycles.
 - Similar to the multi-cycle datapath
- ❑ Design a separate datapath **stage** for the execution performed during each cycle.
 - Build **pipeline registers** to communicate between the stages.

Our new pipelined datapath

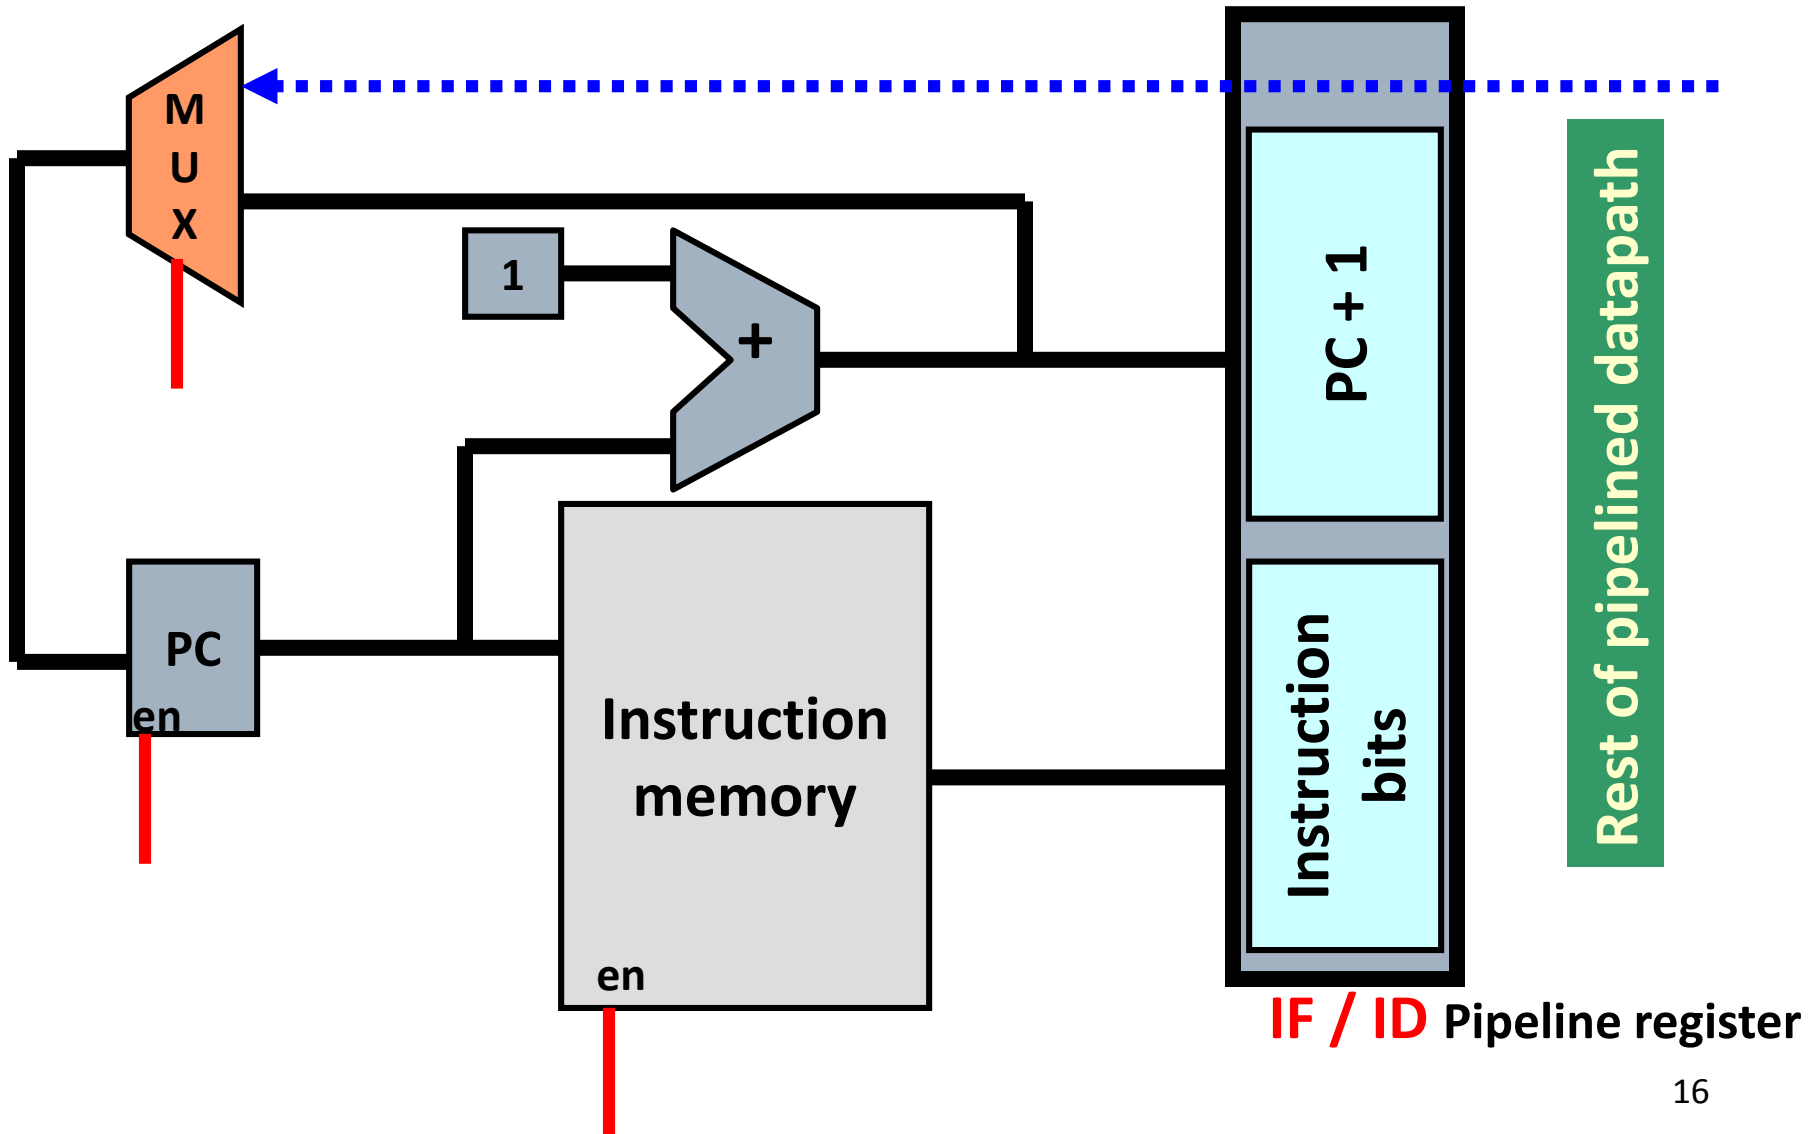


Stage 1: Fetch

- ❑ Design a datapath that can fetch an instruction from memory every cycle.
 - Use PC to index memory to read instruction
 - Increment the PC (assume no branches for now)

- ❑ Write everything needed to complete execution to the **pipeline register (IF/ID)**
 - The next **stage** will read this pipeline register.
 - Note that pipeline register must be edge-triggered

Pipeline datapath – Fetch stage

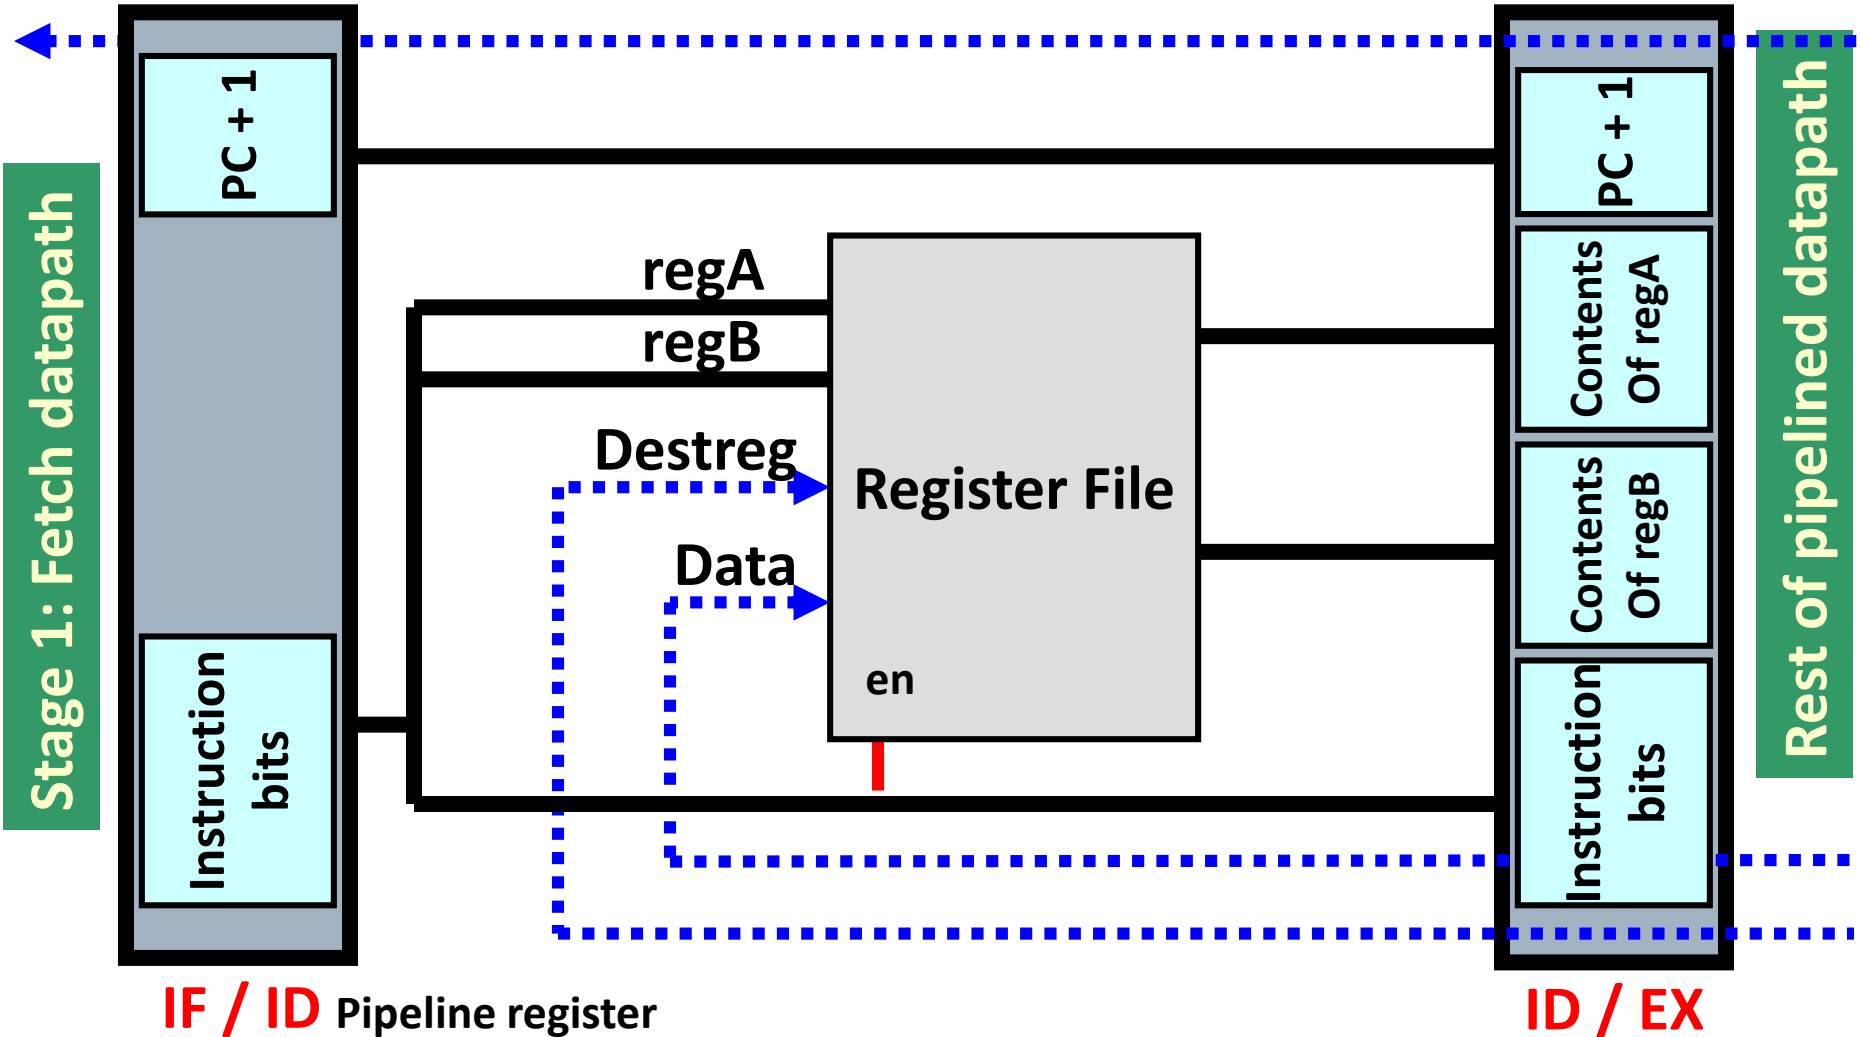


Stage 2: Decode

- ❑ Design a datapath that reads the IF/ID pipeline register, decodes instruction and reads register file (specified by regA and regB of instruction bits).
 - Decode is easy, just pass on the opcode and let later stages figure out their own control signals for the instruction.

- ❑ Write everything needed to complete execution to the **pipeline register (ID/EX)**
 - Pass on the offset field and both destination register specifiers (or simply pass on the whole instruction!).
 - Including PC+1 even though decode didn't use it.

Pipeline datapath – Decode stage

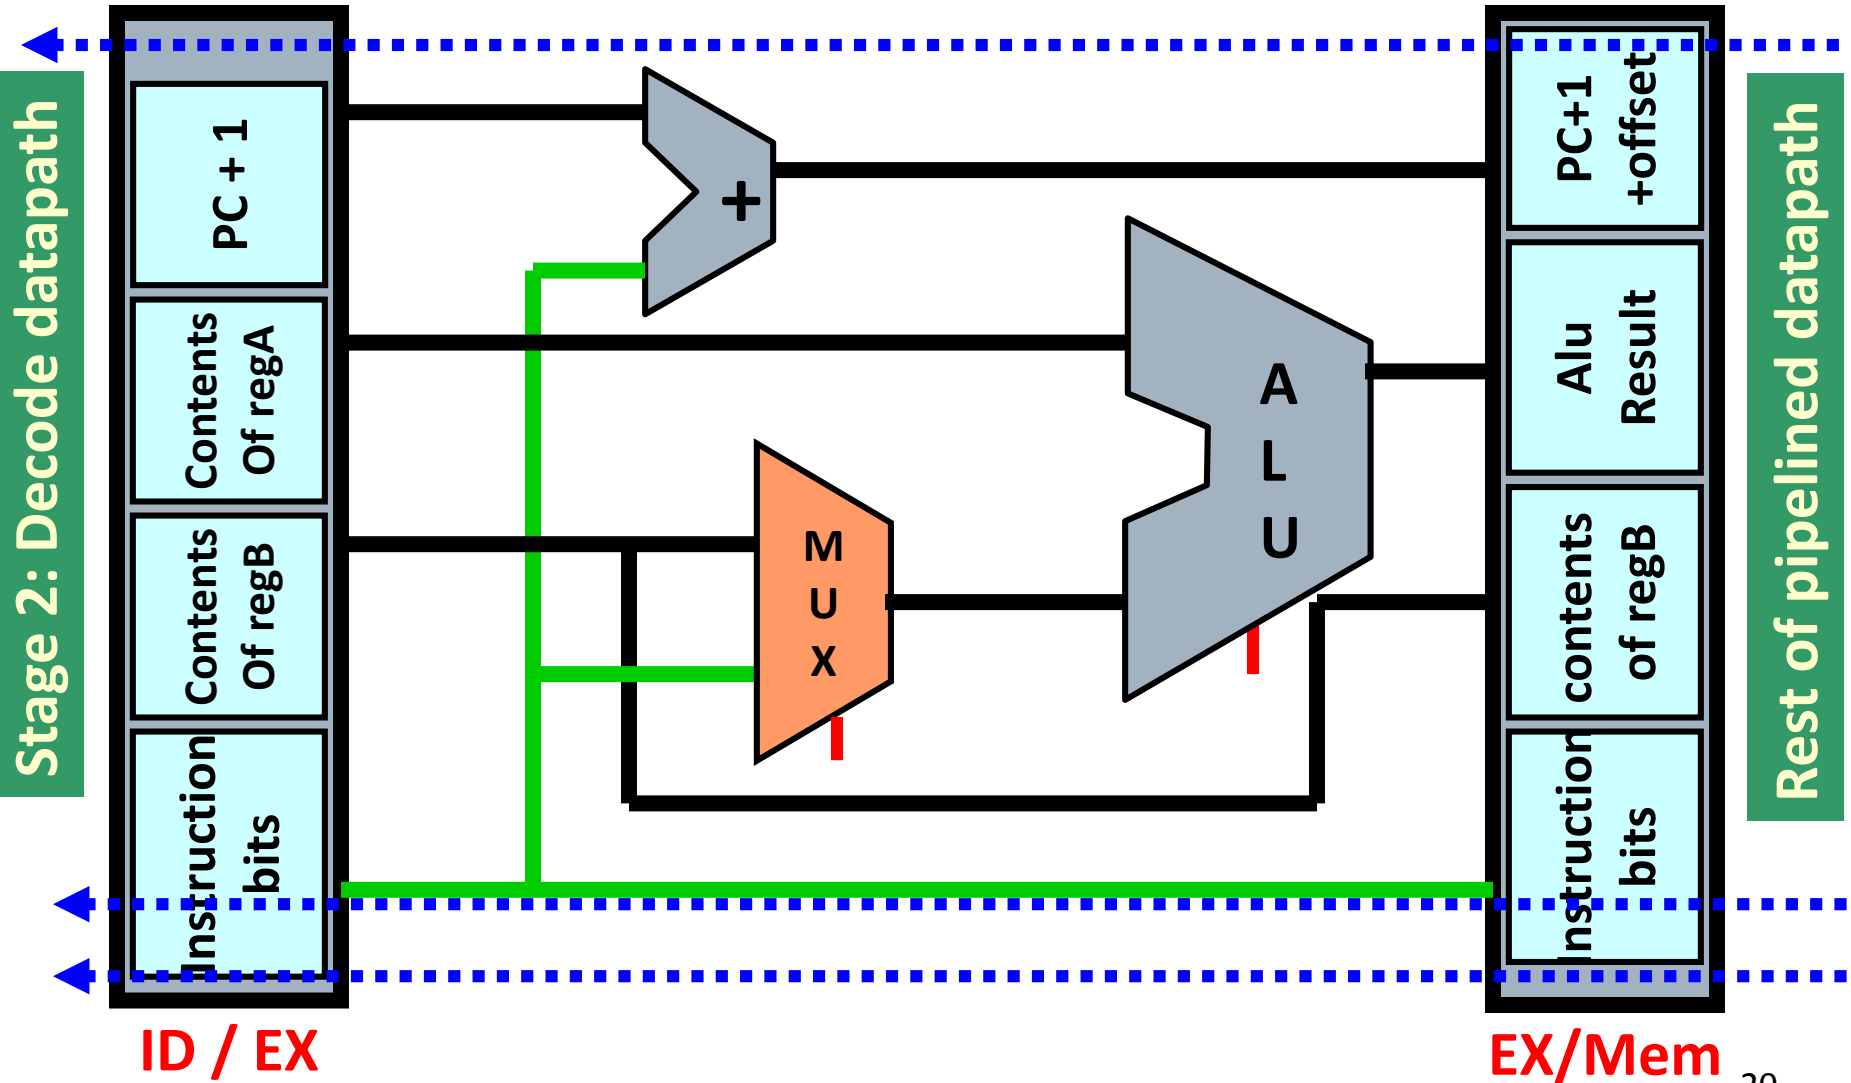


Stage 3: Execute

- ❑ Design a datapath that performs the proper ALU operation for the instruction specified and the values present in the ID/EX pipeline register.
 - The inputs are the contents of regA and either the contents of regB or the offset field on the instruction.
 - Also, calculate $PC+1+offset$ in case this is a branch.

- ❑ Write everything needed to complete execution to the pipeline register (EX/Mem)
 - ALU result, contents of regB and $PC+1+offset$
 - Instruction bits for opcode and destReg specifiers
 - Result from comparison of regA and regB contents

Pipeline datapath – Execute stage

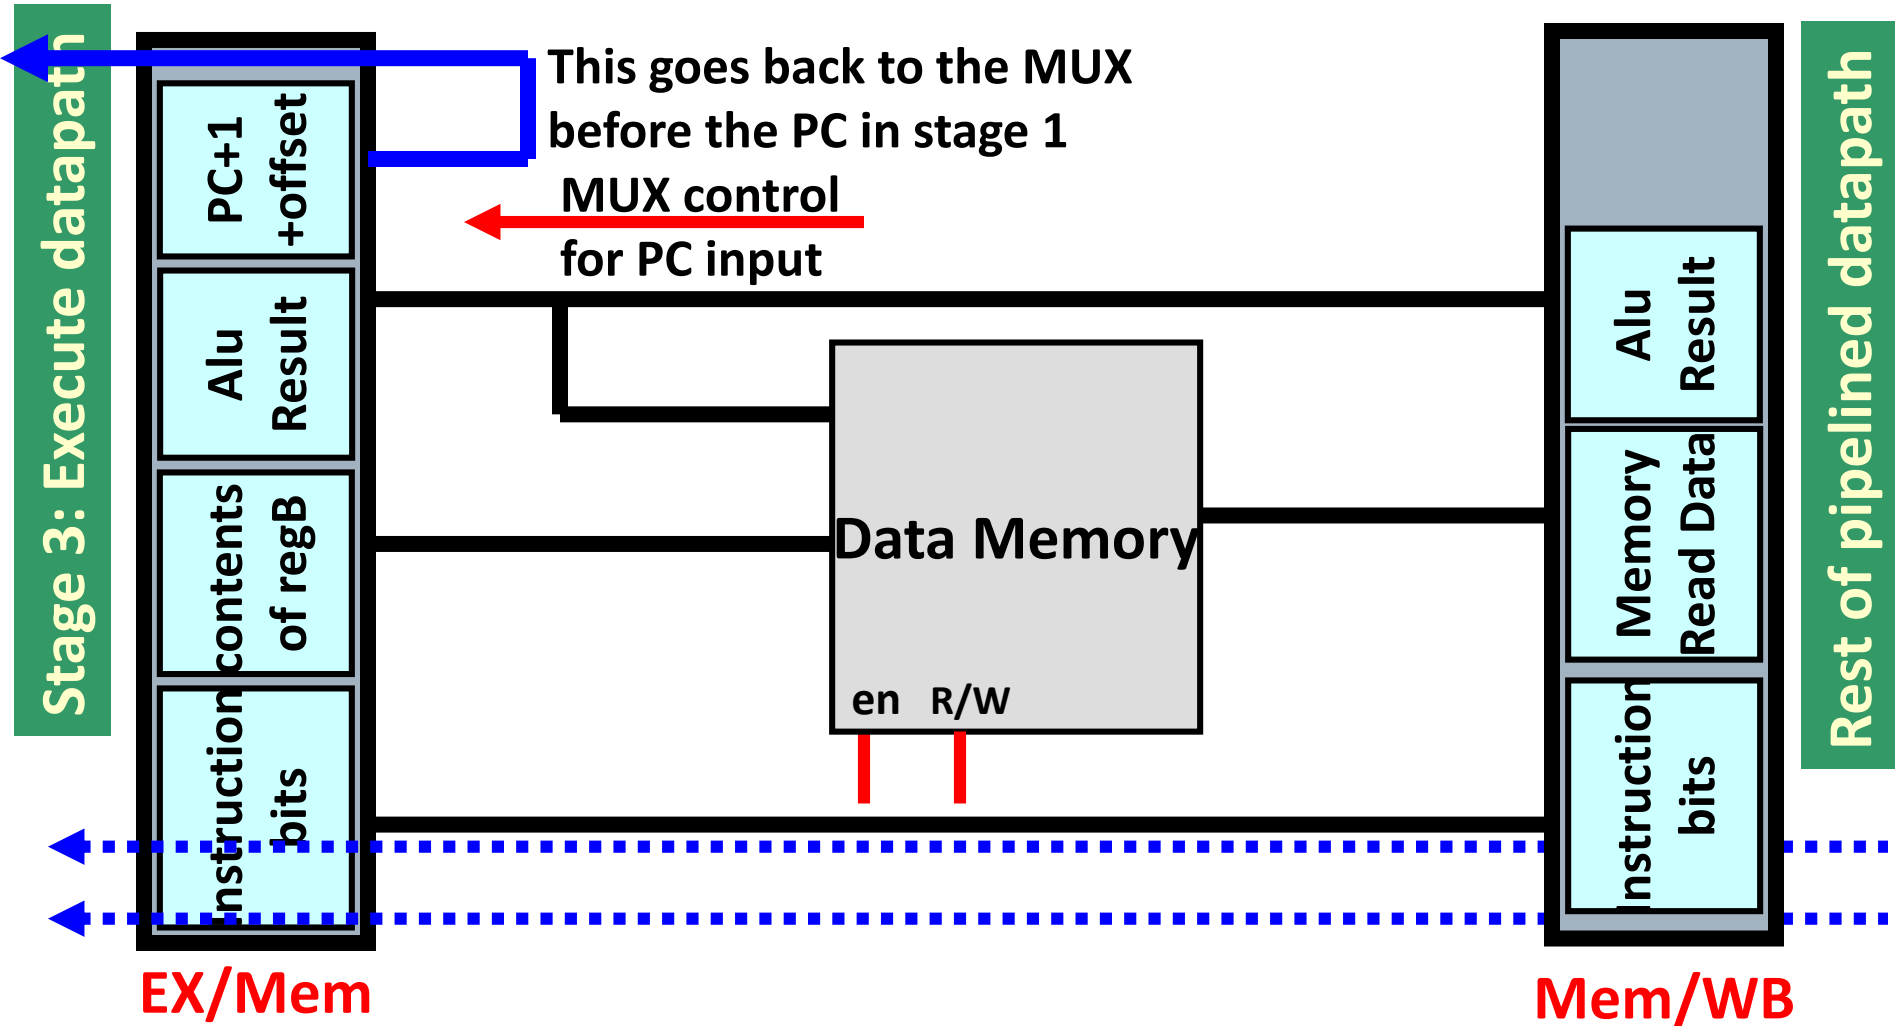


Stage 4: Memory Operation

- ❑ Design a datapath that performs the proper memory operation for the instruction specified and the values present in the EX/Mem pipeline register.
 - ALU result contains address for **ld** and **st** instructions.
 - Opcode bits control memory R/W and enable signals.

- ❑ Write everything needed to complete execution to the **pipeline register (Mem/WB)**
 - ALU result and MemData
 - Instruction bits for opcode and destReg specifiers

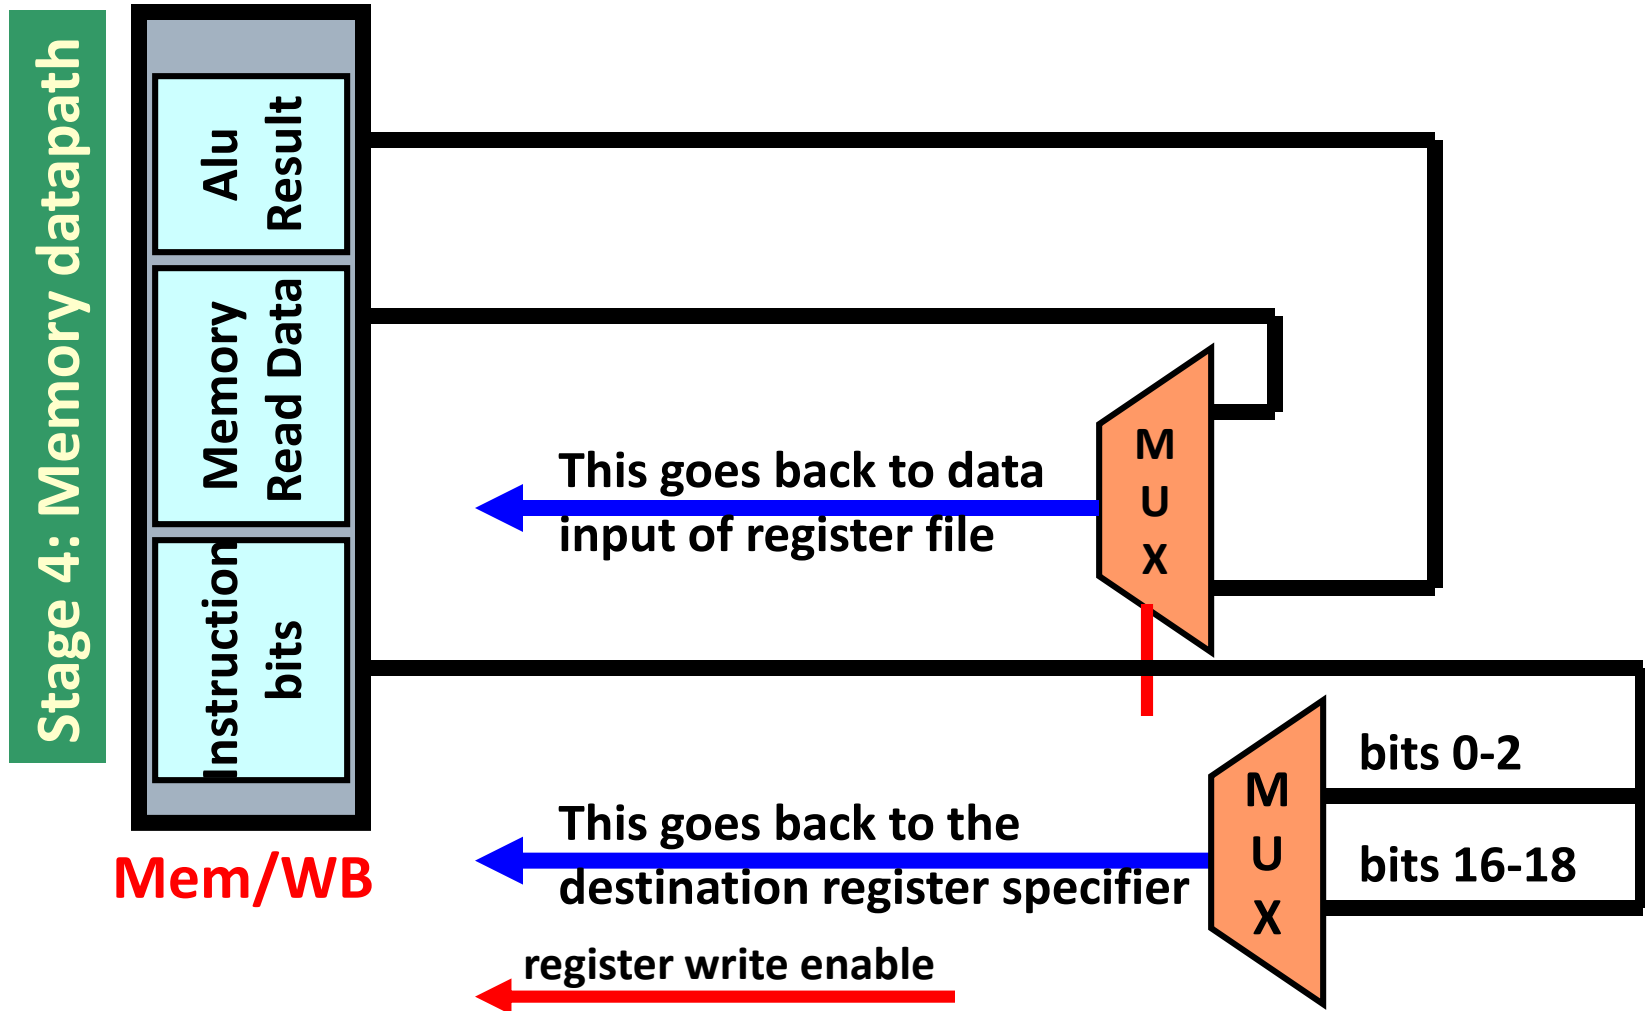
Pipeline datapath – Memory stage



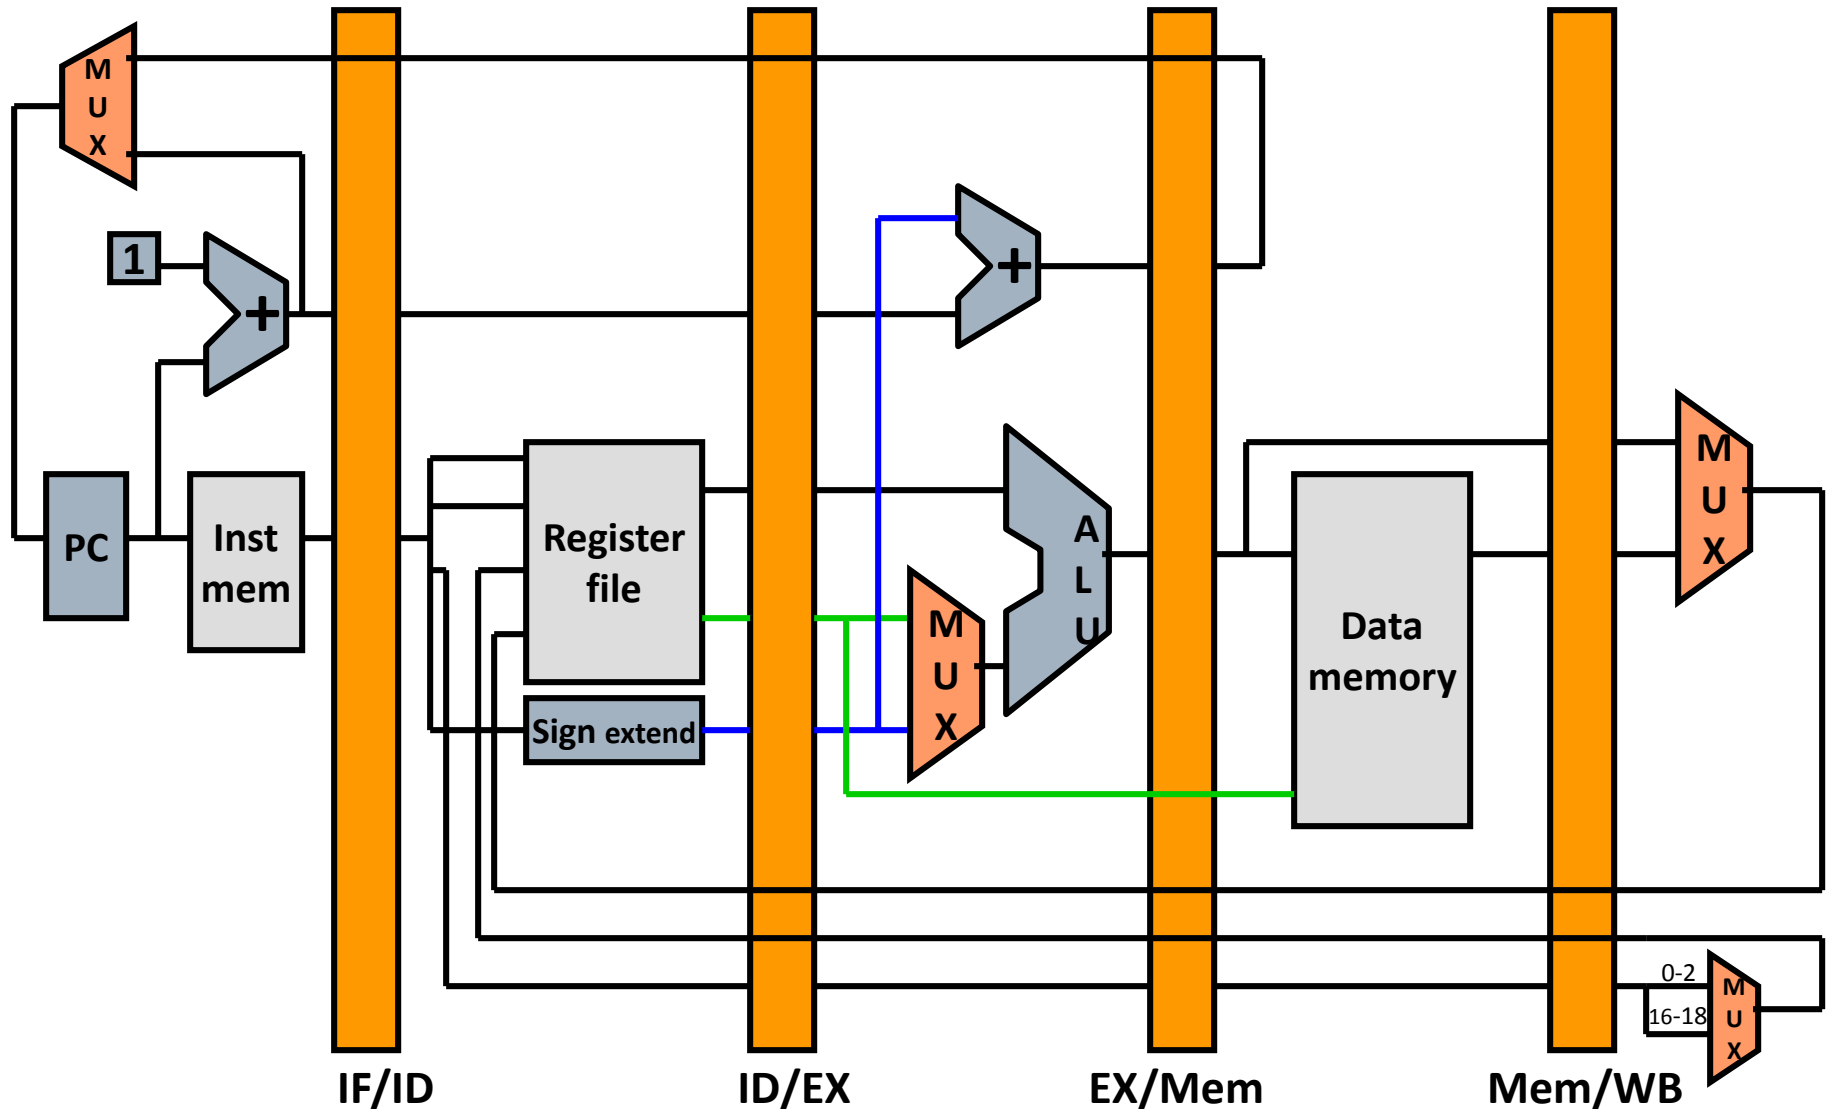
Stage 5: Write back

- ❑ Design a datapath that completes the execution of this instruction, writing to the register file if required.
 - Write MemData to destReg for ld instruction
 - Write ALU result to destReg for add or nand instructions.
 - Opcode bits also control register write enable signal.

Pipeline datapath – Writeback stage



Putting all together

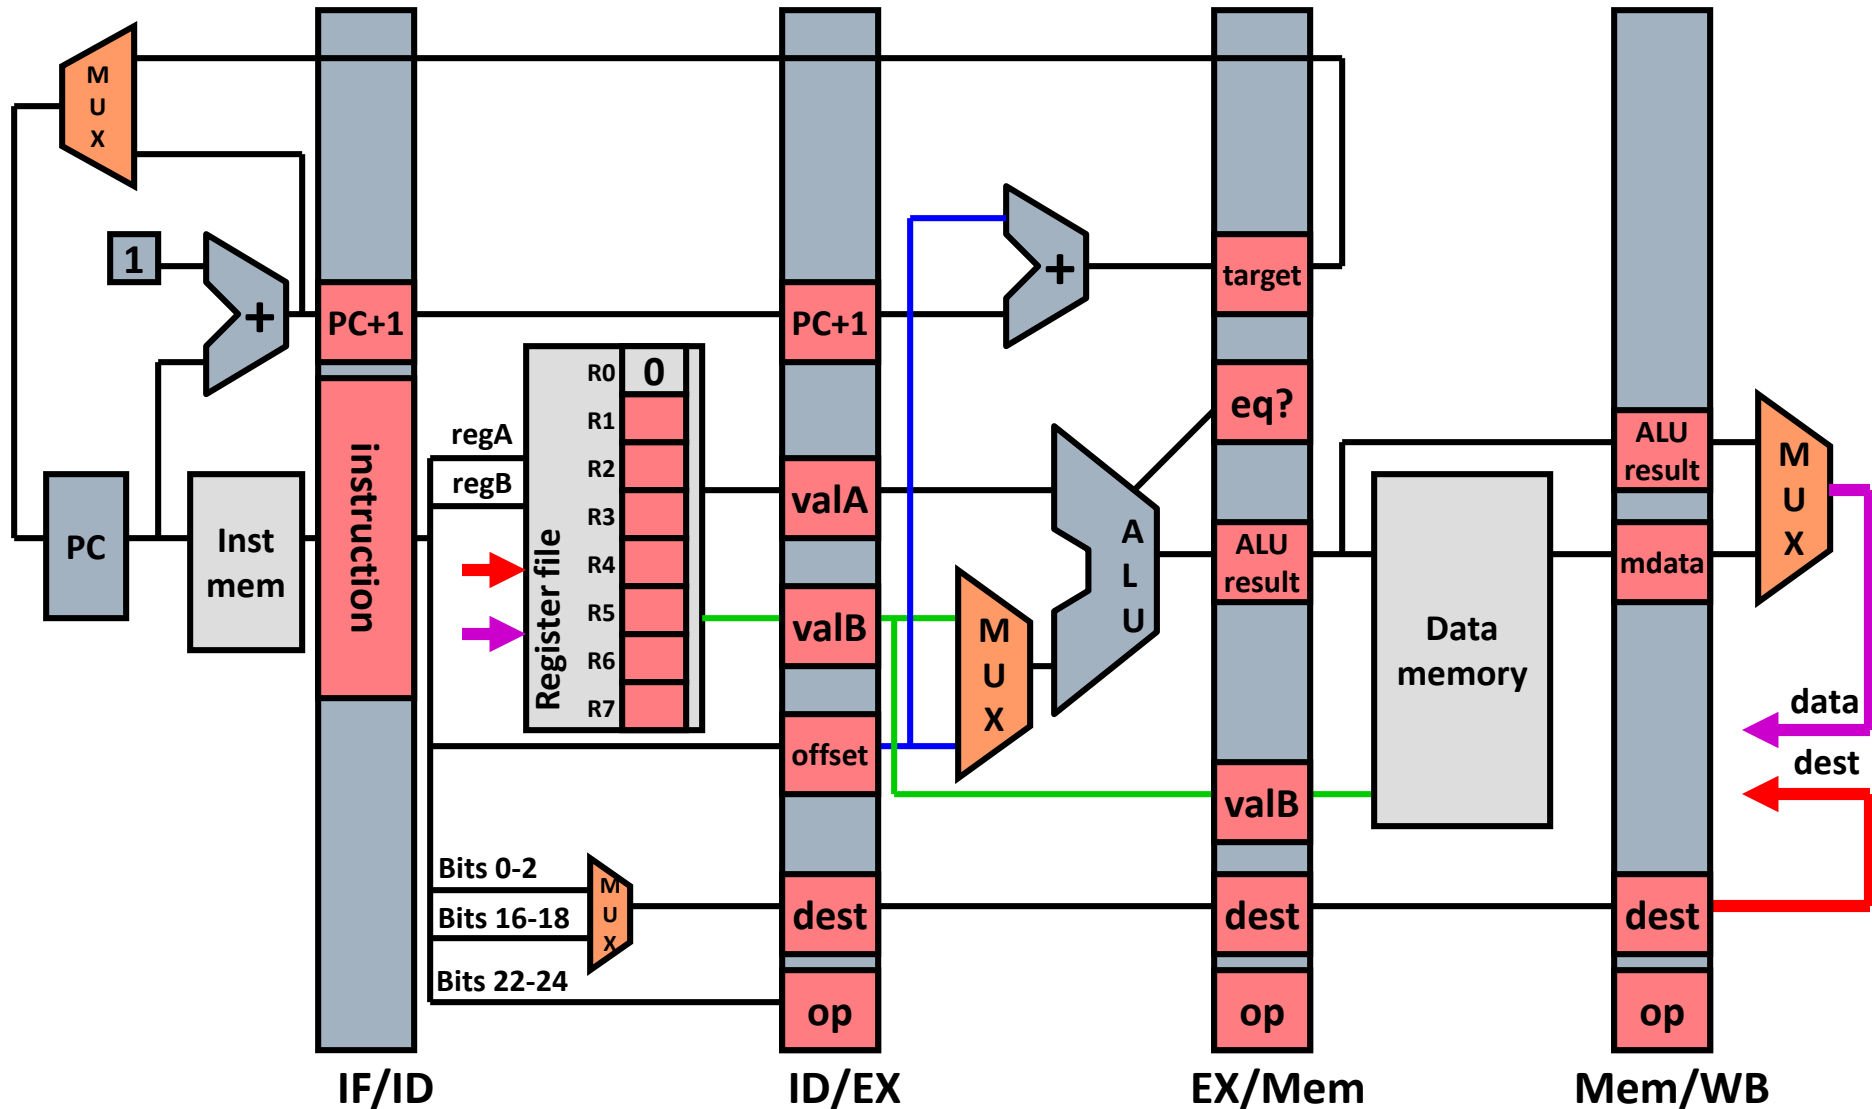


Sample Code (Simple)

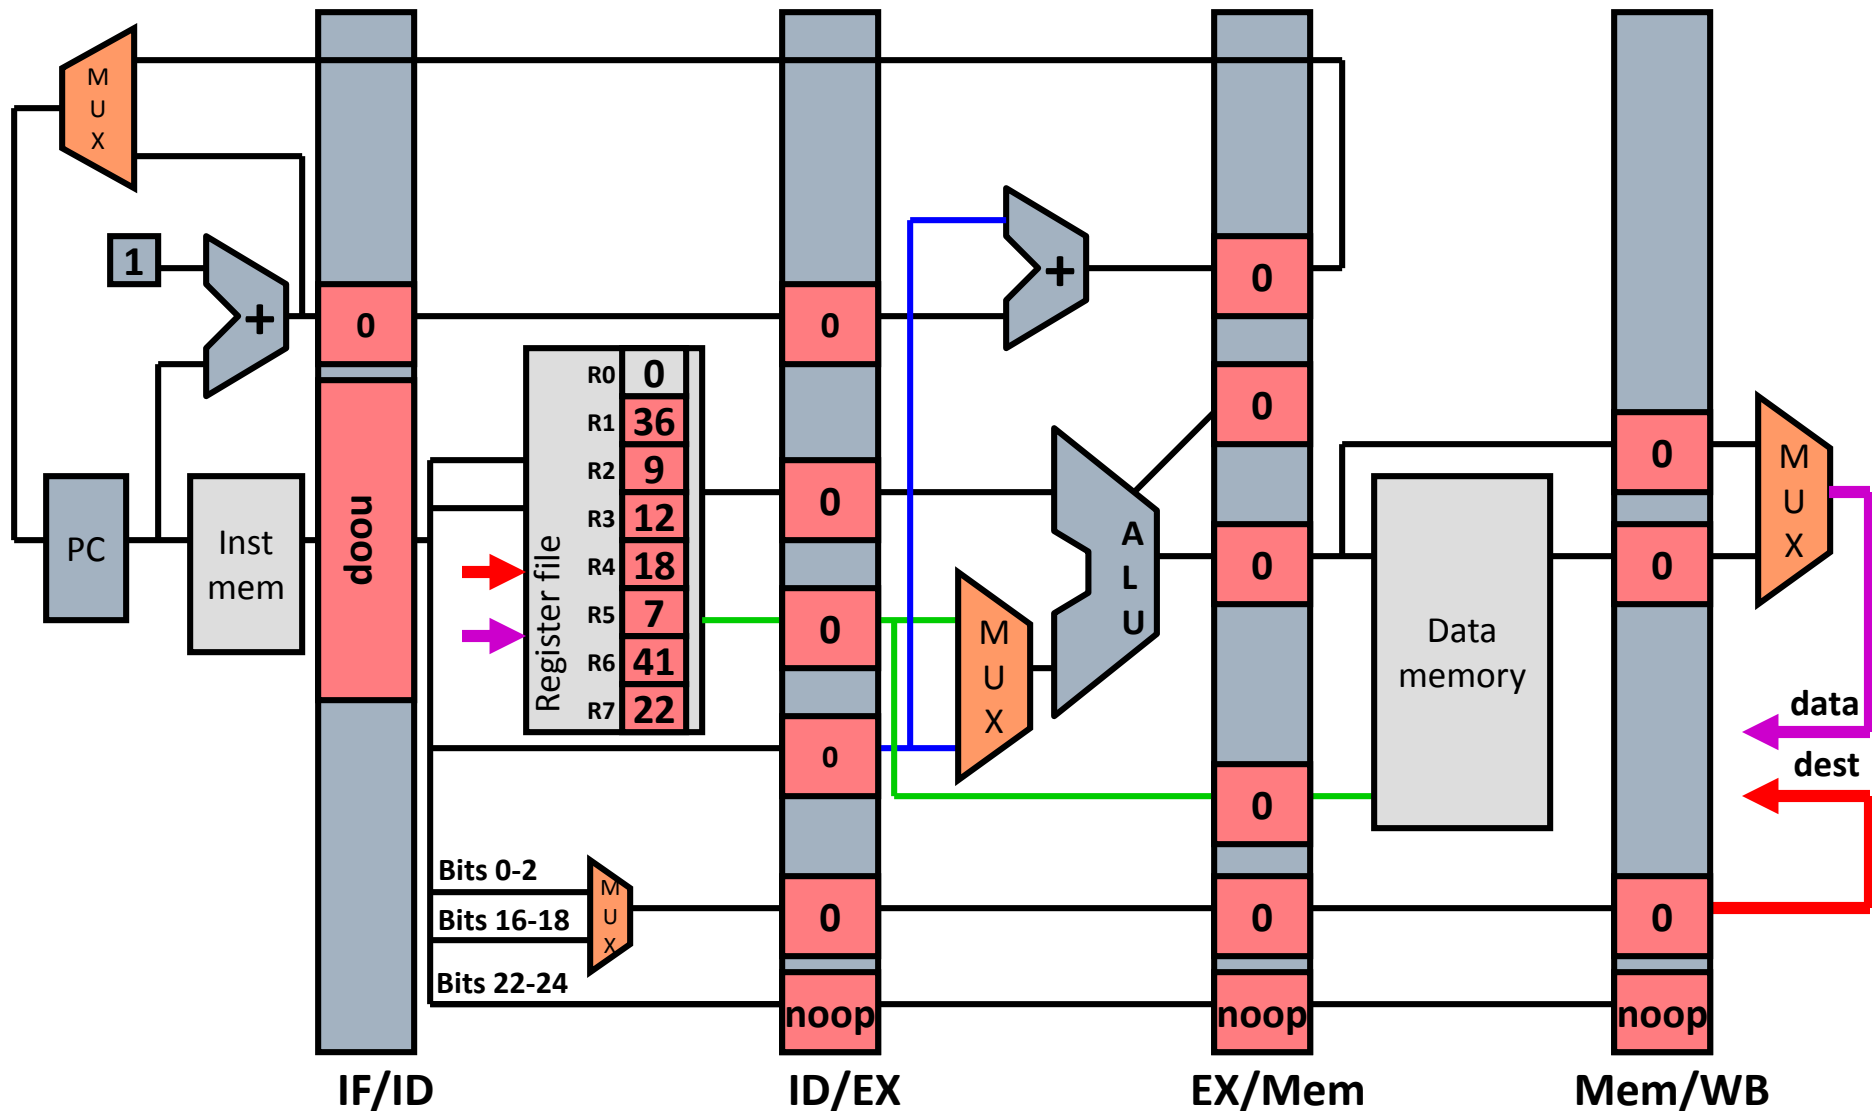
Let's run the following code on pipelined LC2K2x:

- `add 1 2 3 ; reg 3 = reg 1 + reg 2`
- `nand 4 5 6 ; reg 6 = reg 4 & reg 5`
- `lw 2 4 20 ; reg 4 = Mem[reg2+20]`
- `add 2 5 5 ; reg 5 = reg 2 + reg 5`
- `sw 3 7 10 ; Mem[reg3+10] = reg 7`

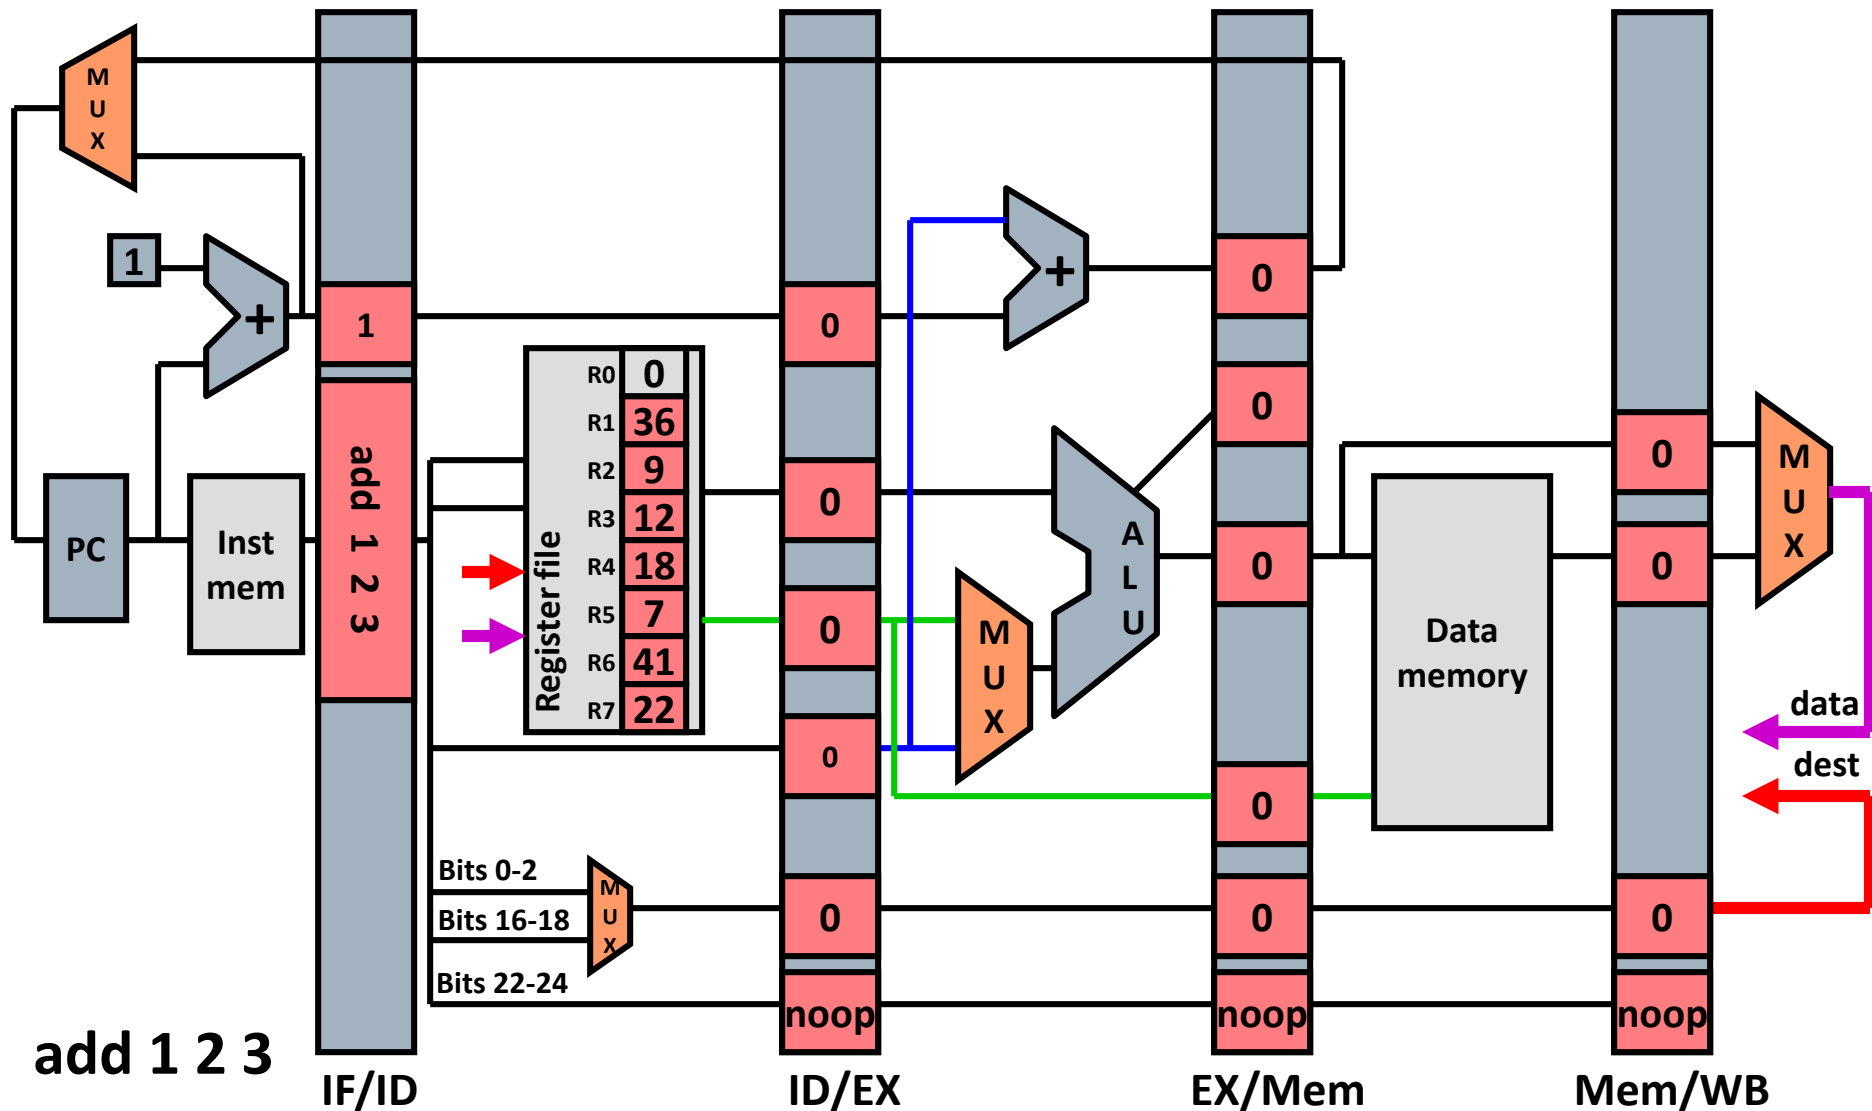
Pipeline datapath



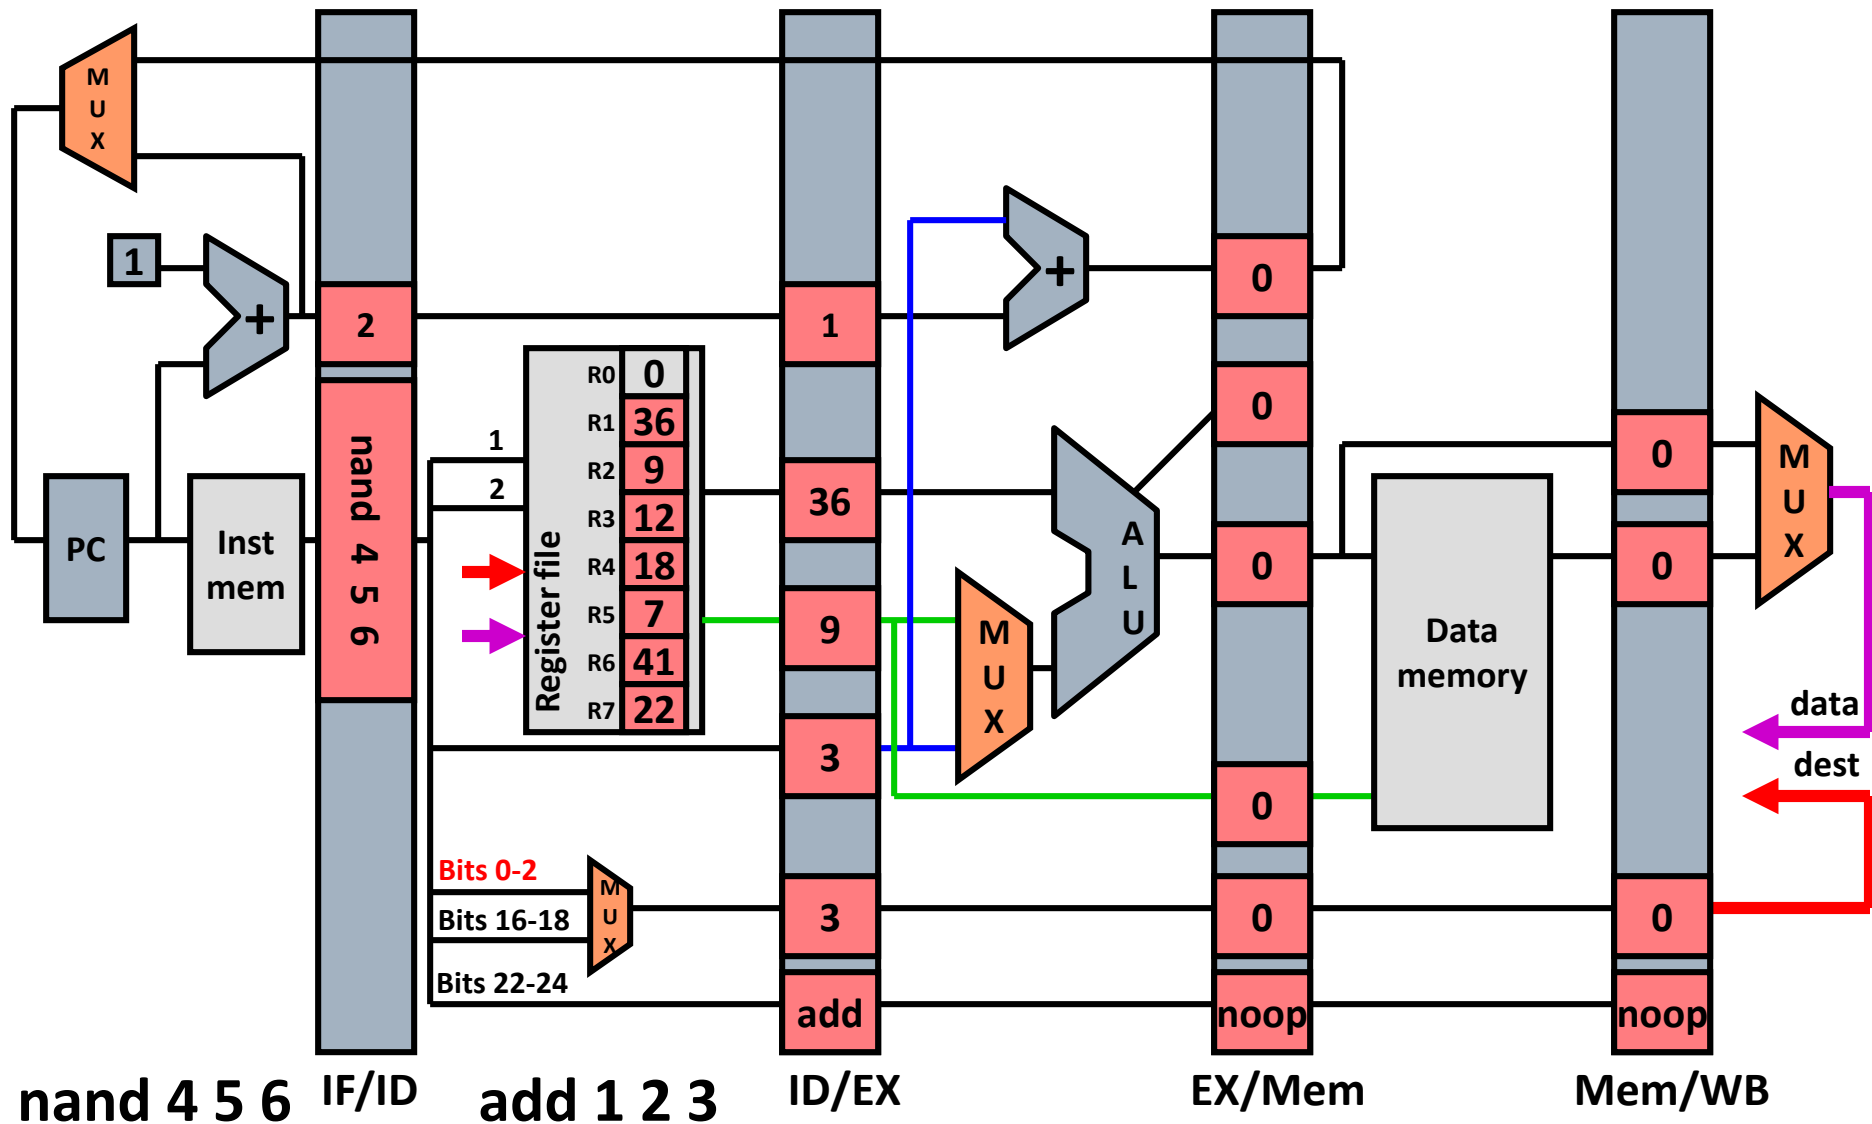
Time 0 - Initial state



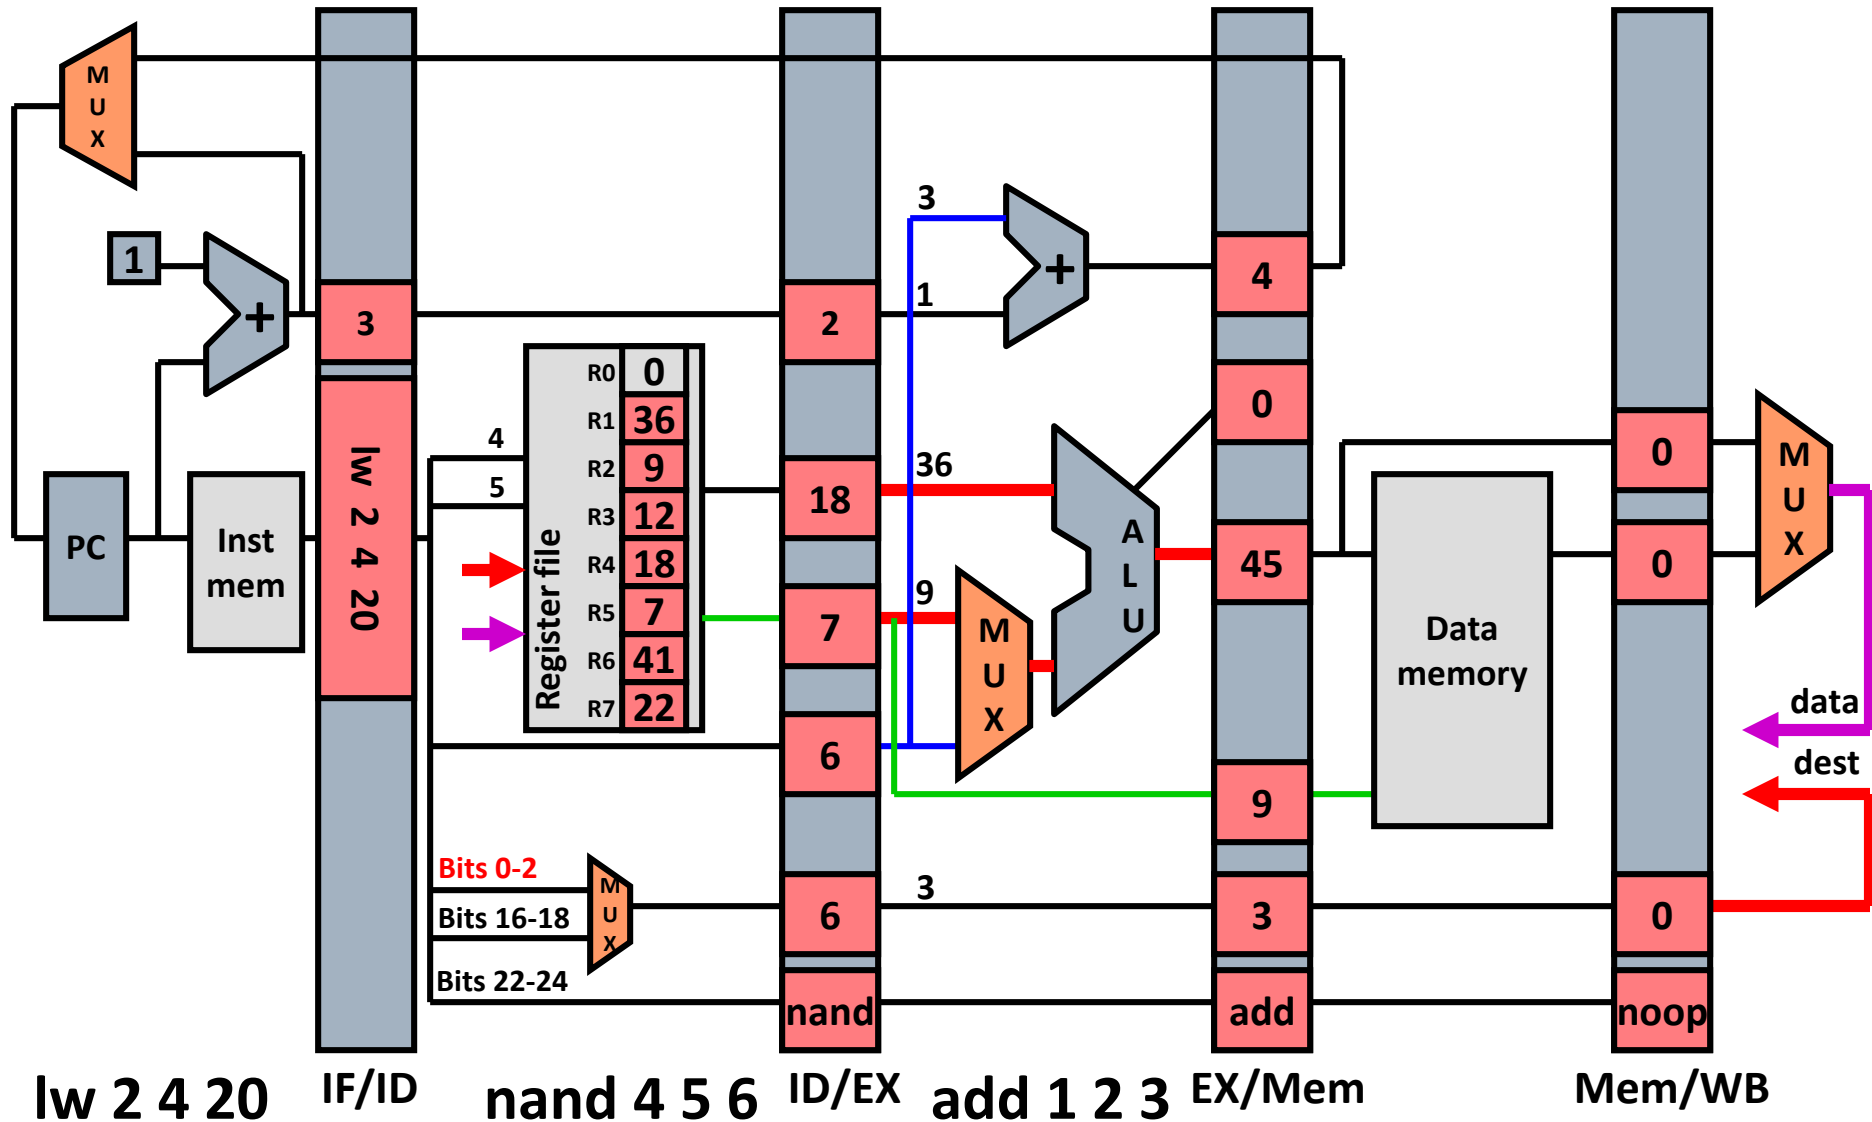
Time 1 - Fetch: add 1 2 3



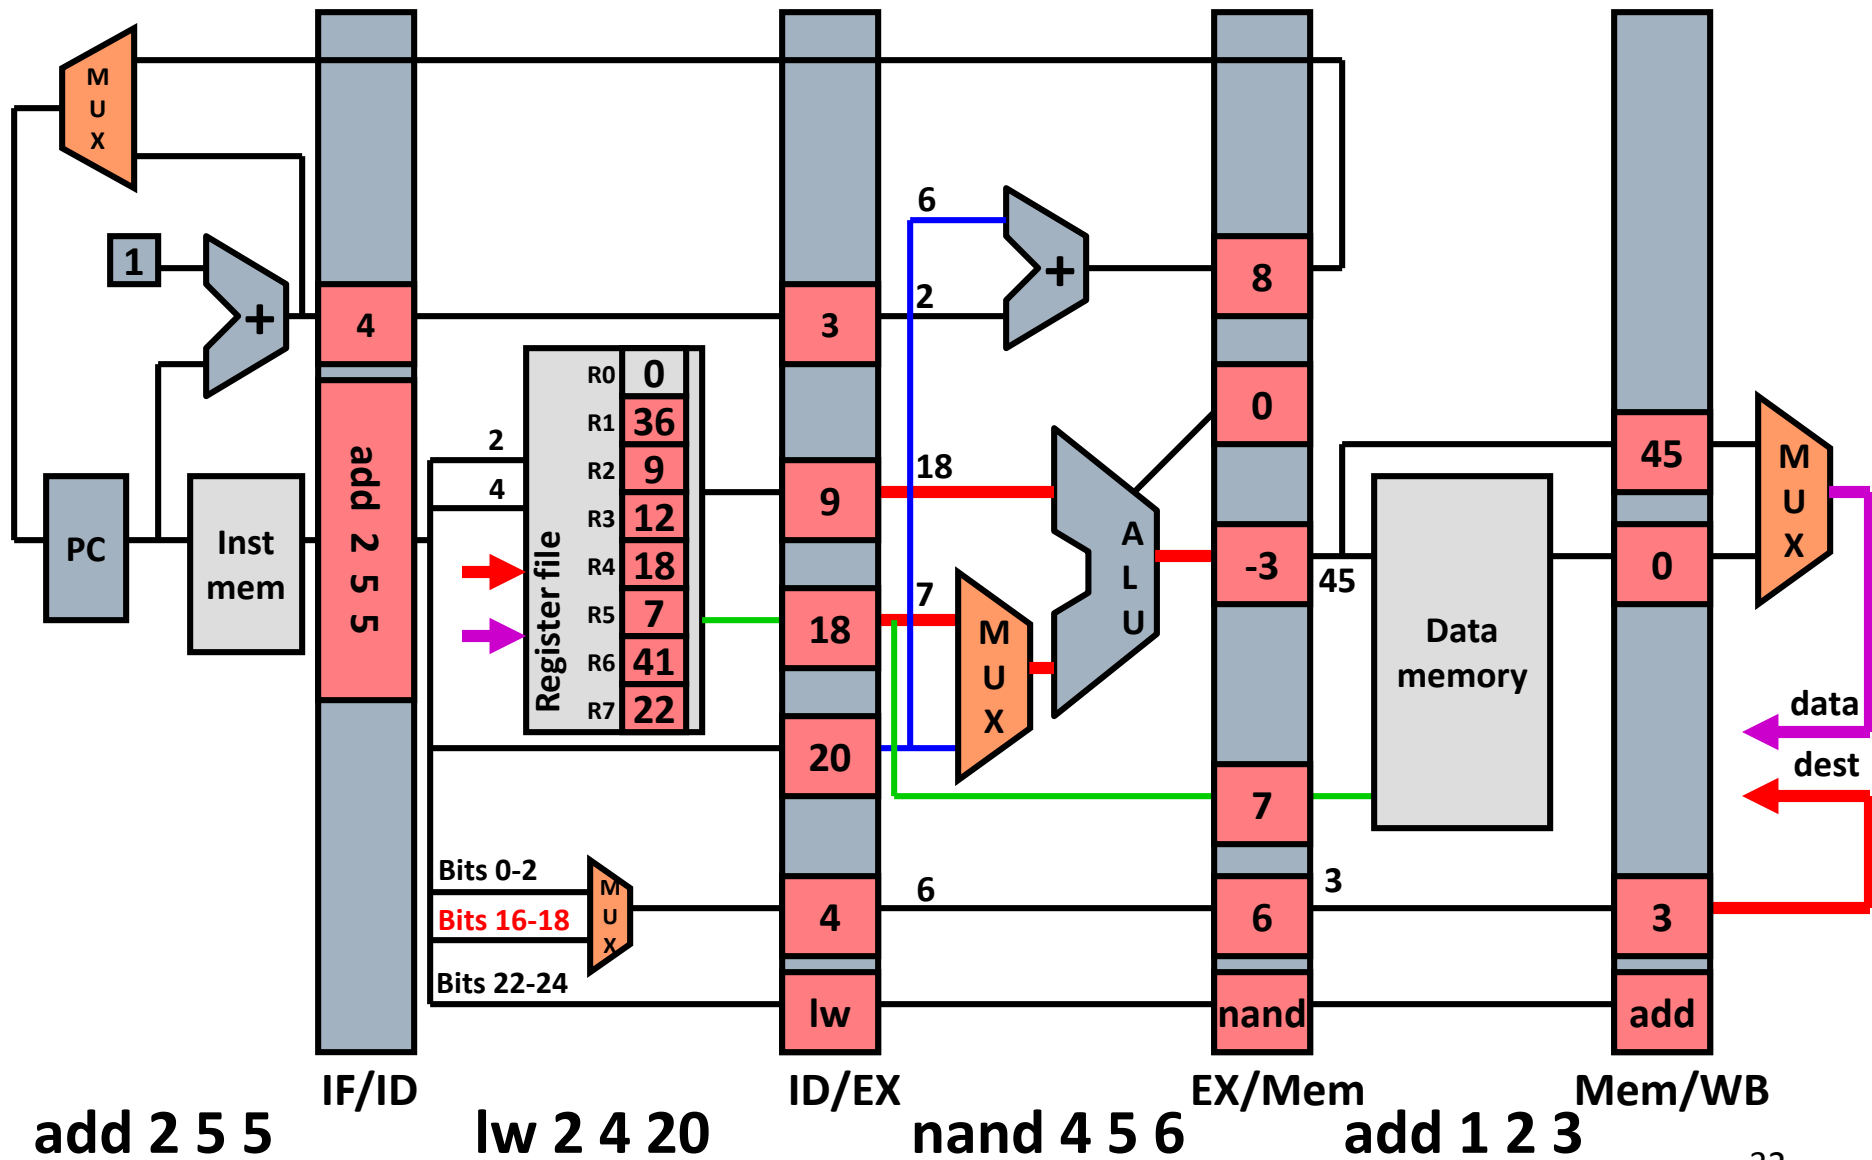
Time 2 - Fetch: nand 4 5 6



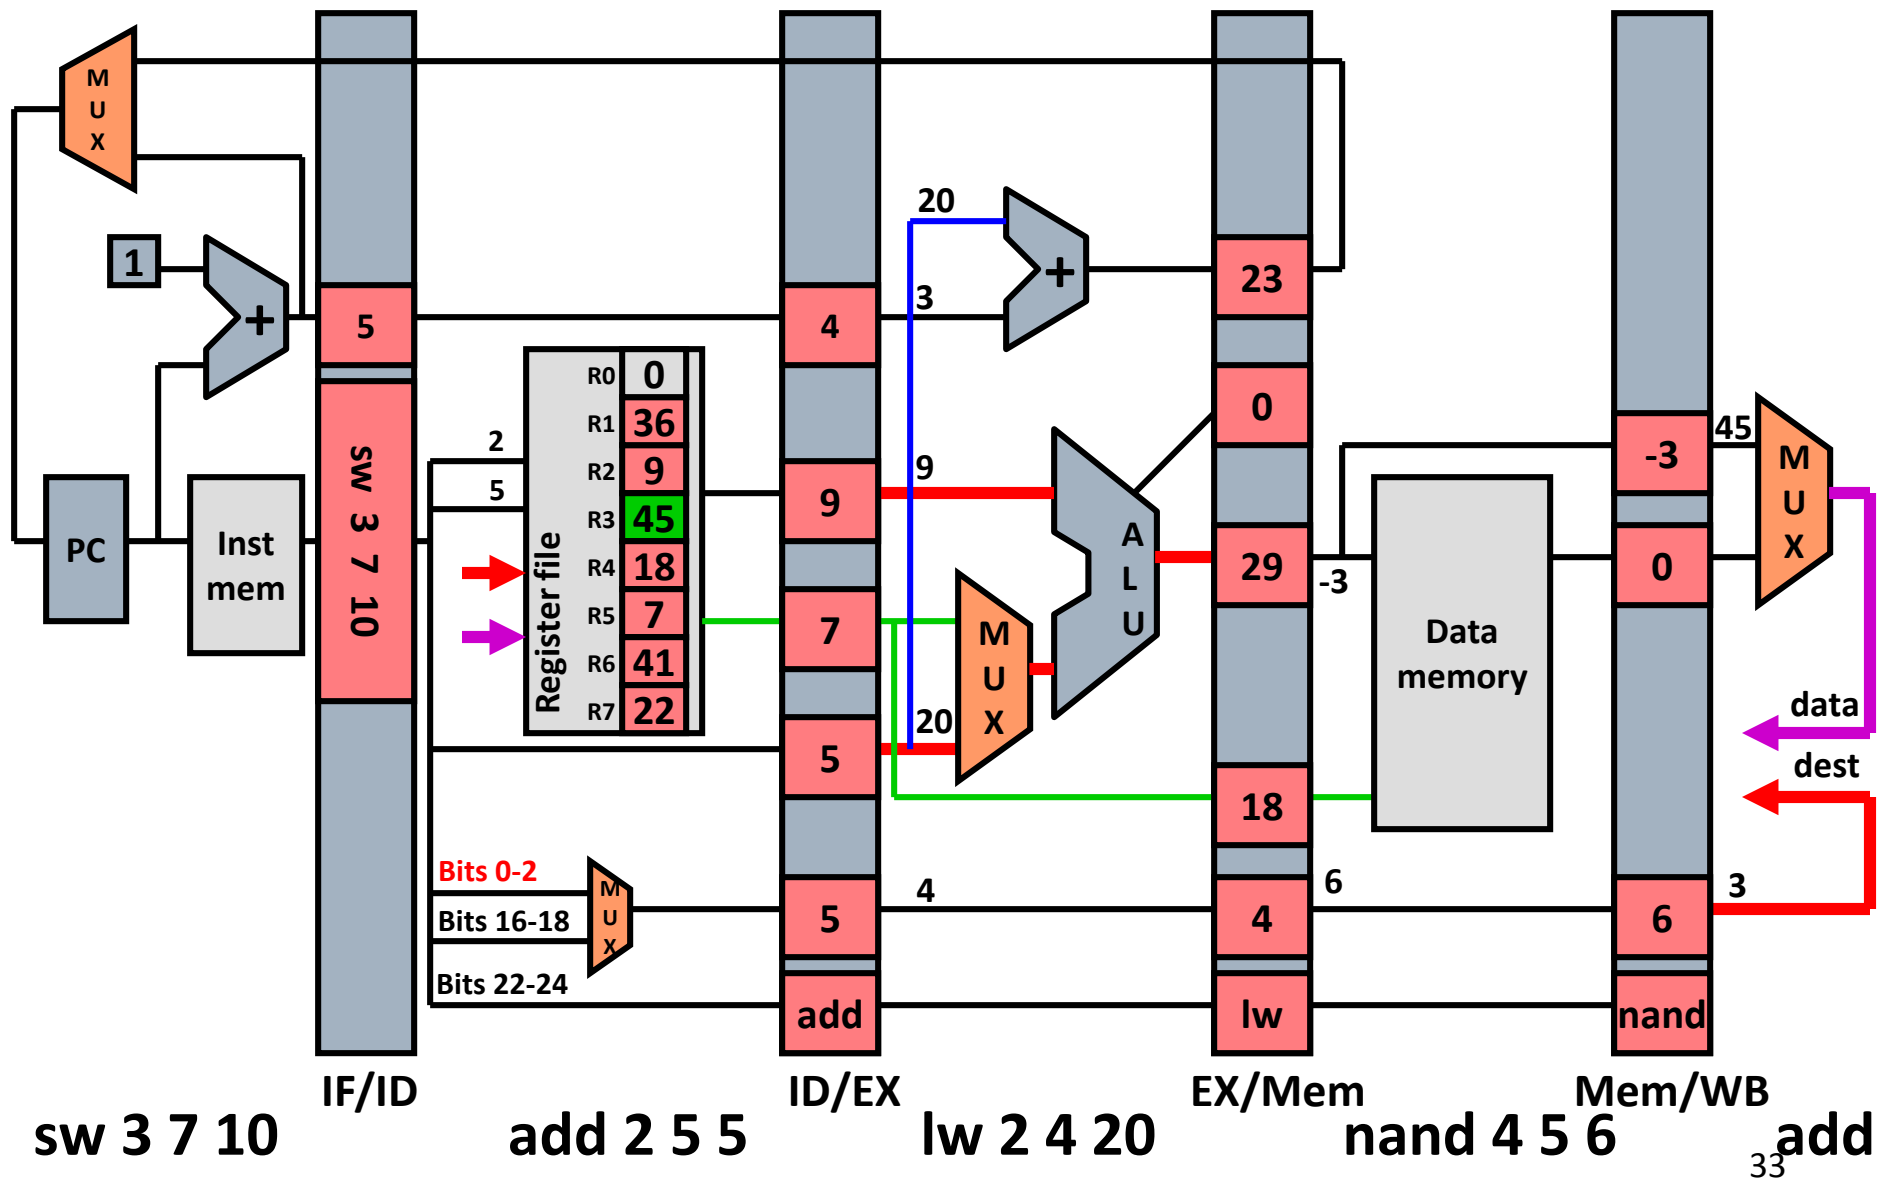
Time 3 - Fetch: lw 2 4 20



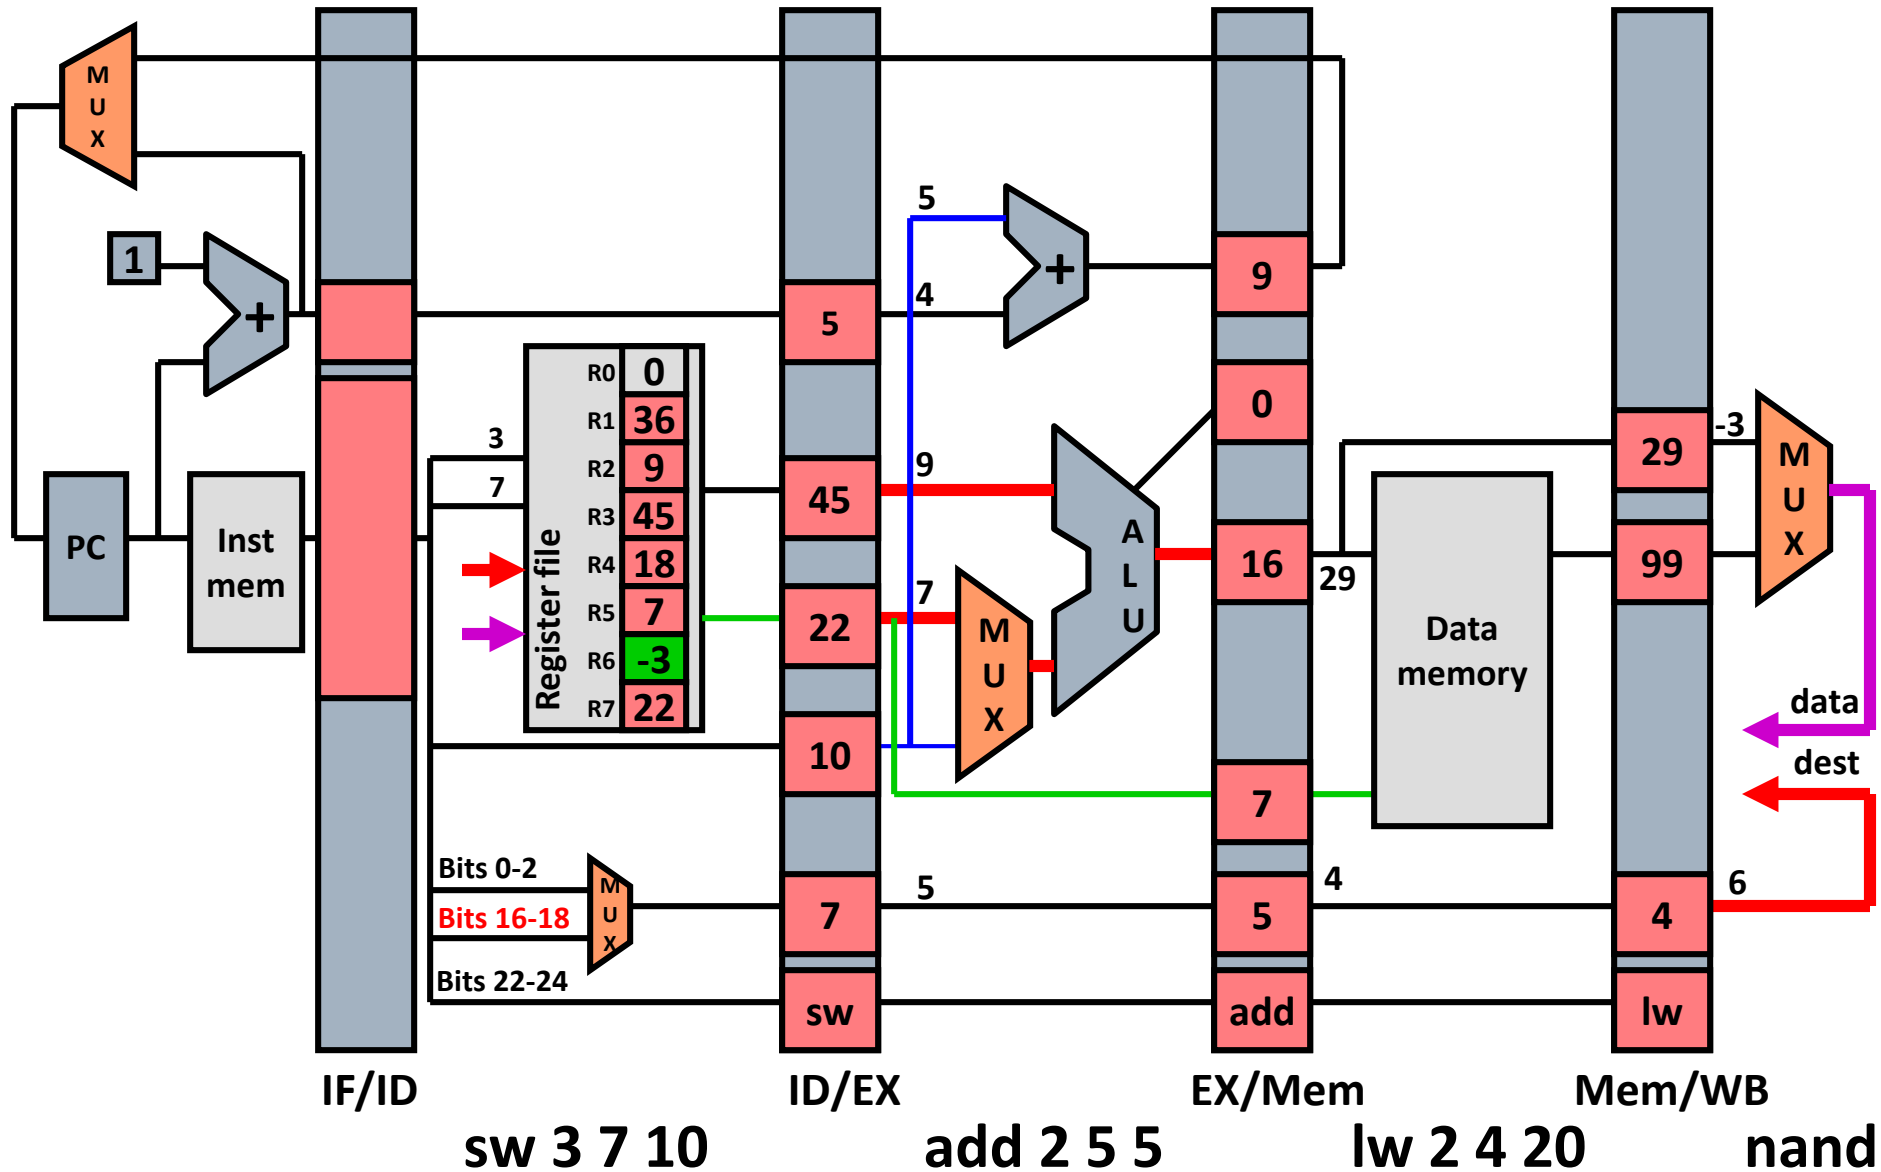
Time 4 - Fetch: add 2 5 5



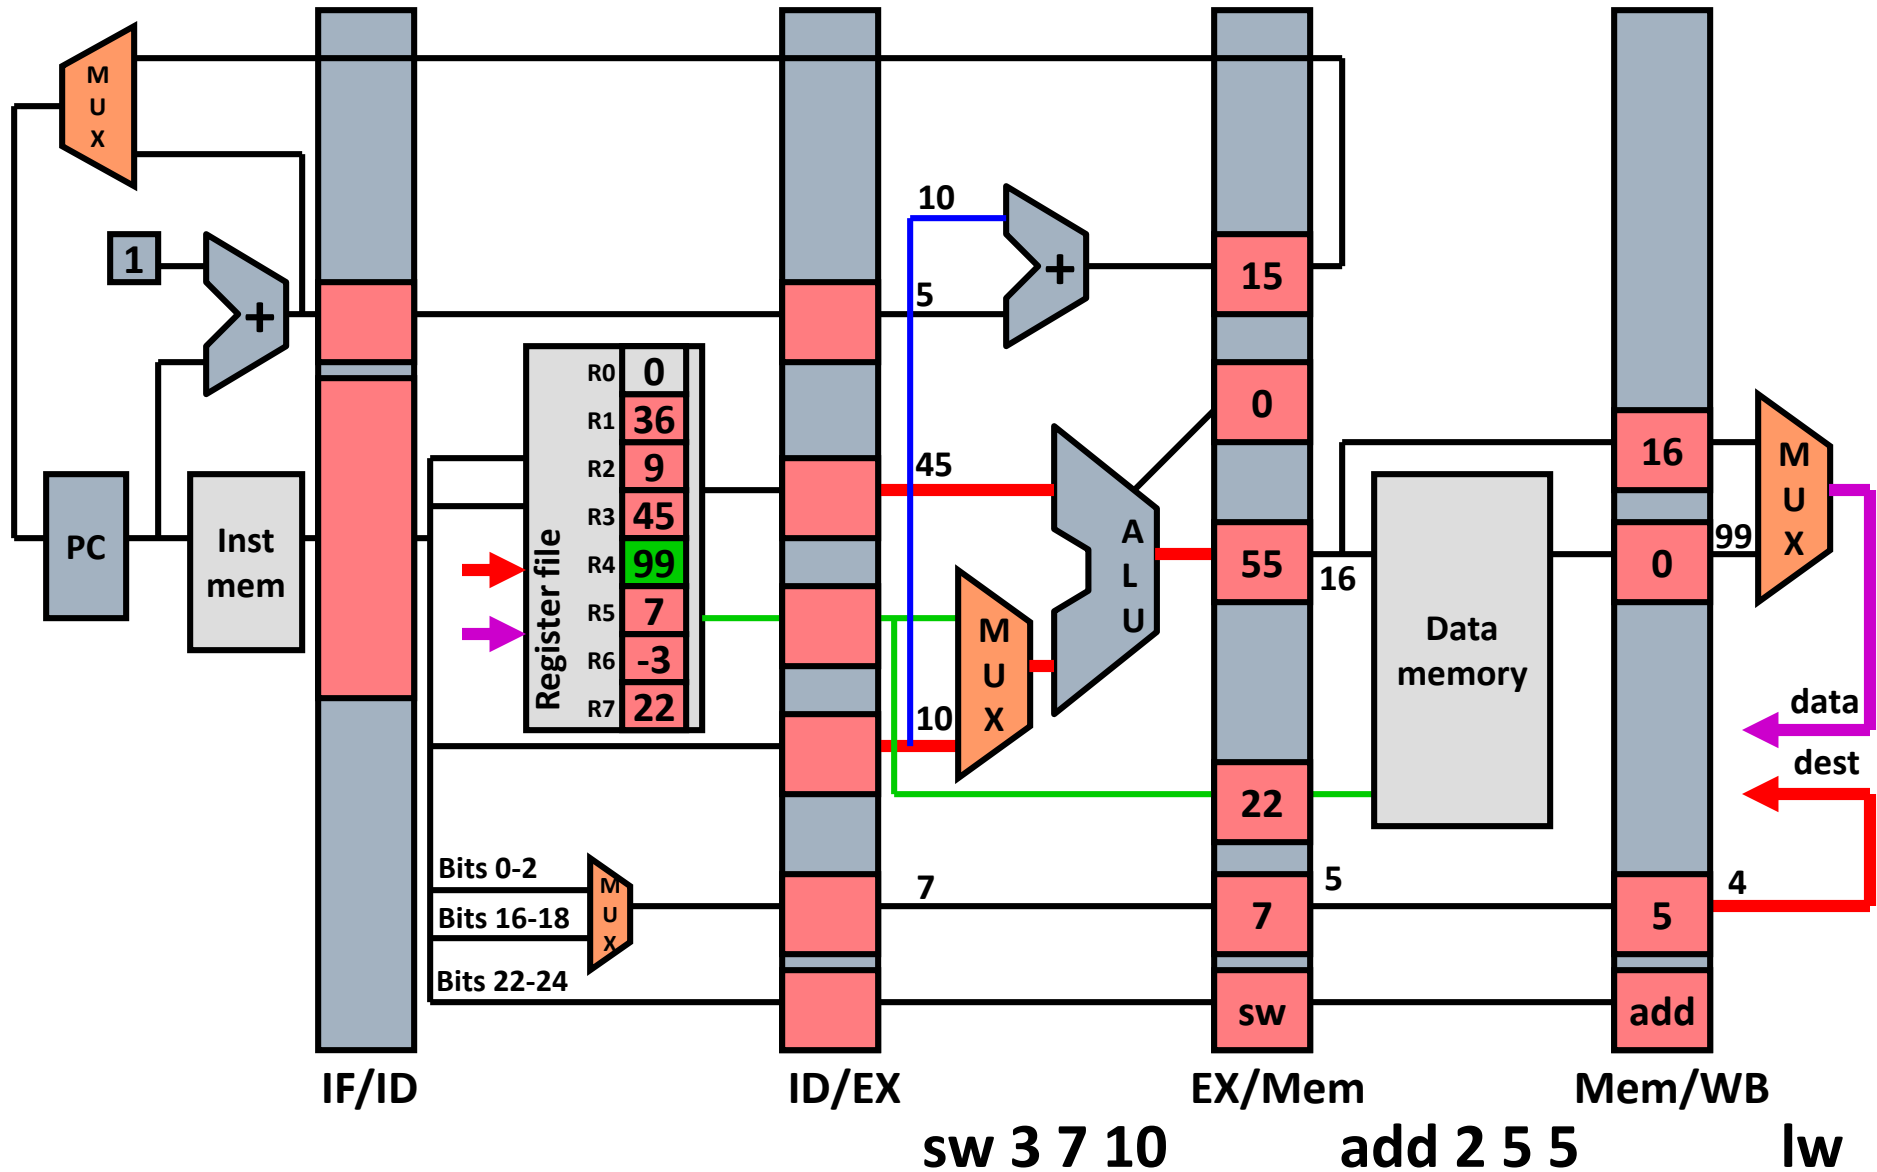
Time 5 - Fetch: sw 3 7 10



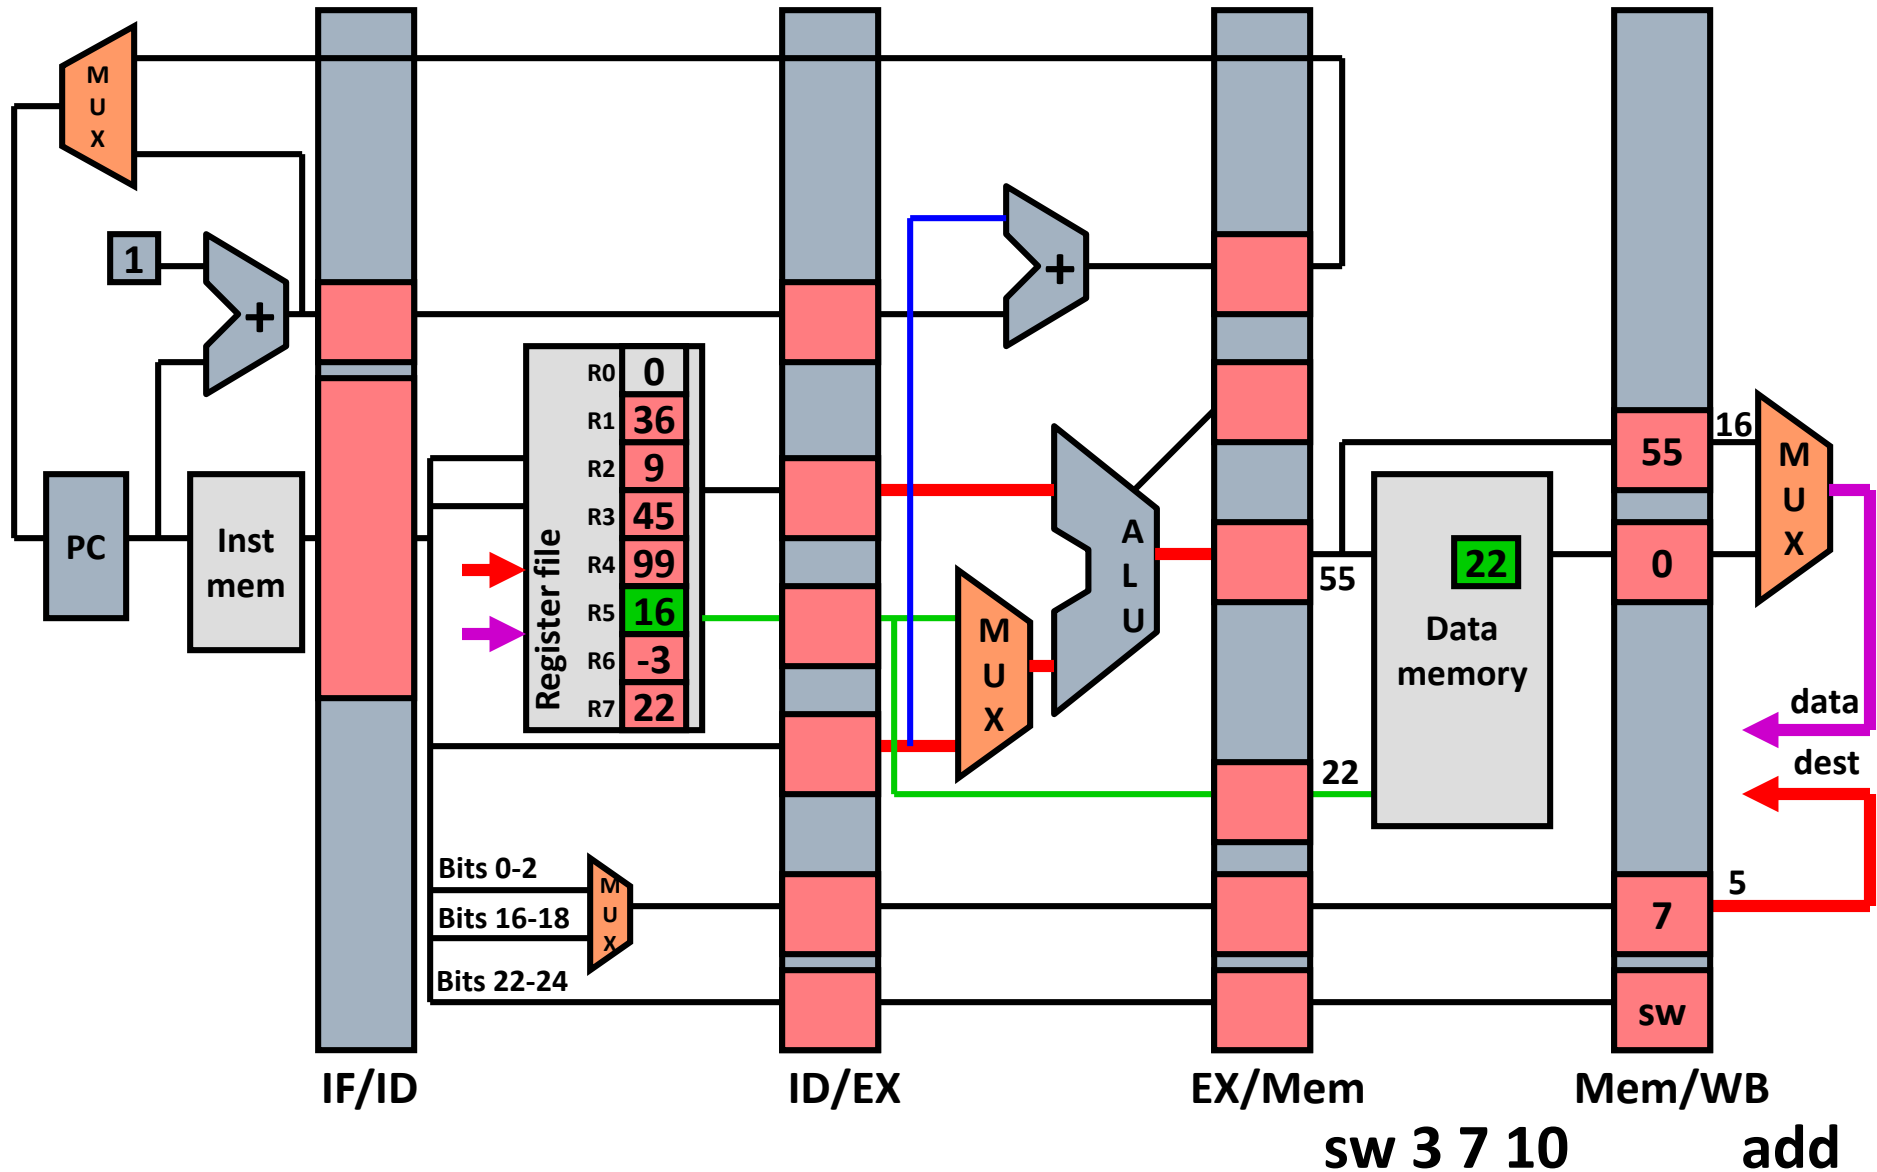
Time 6 – no more instructions



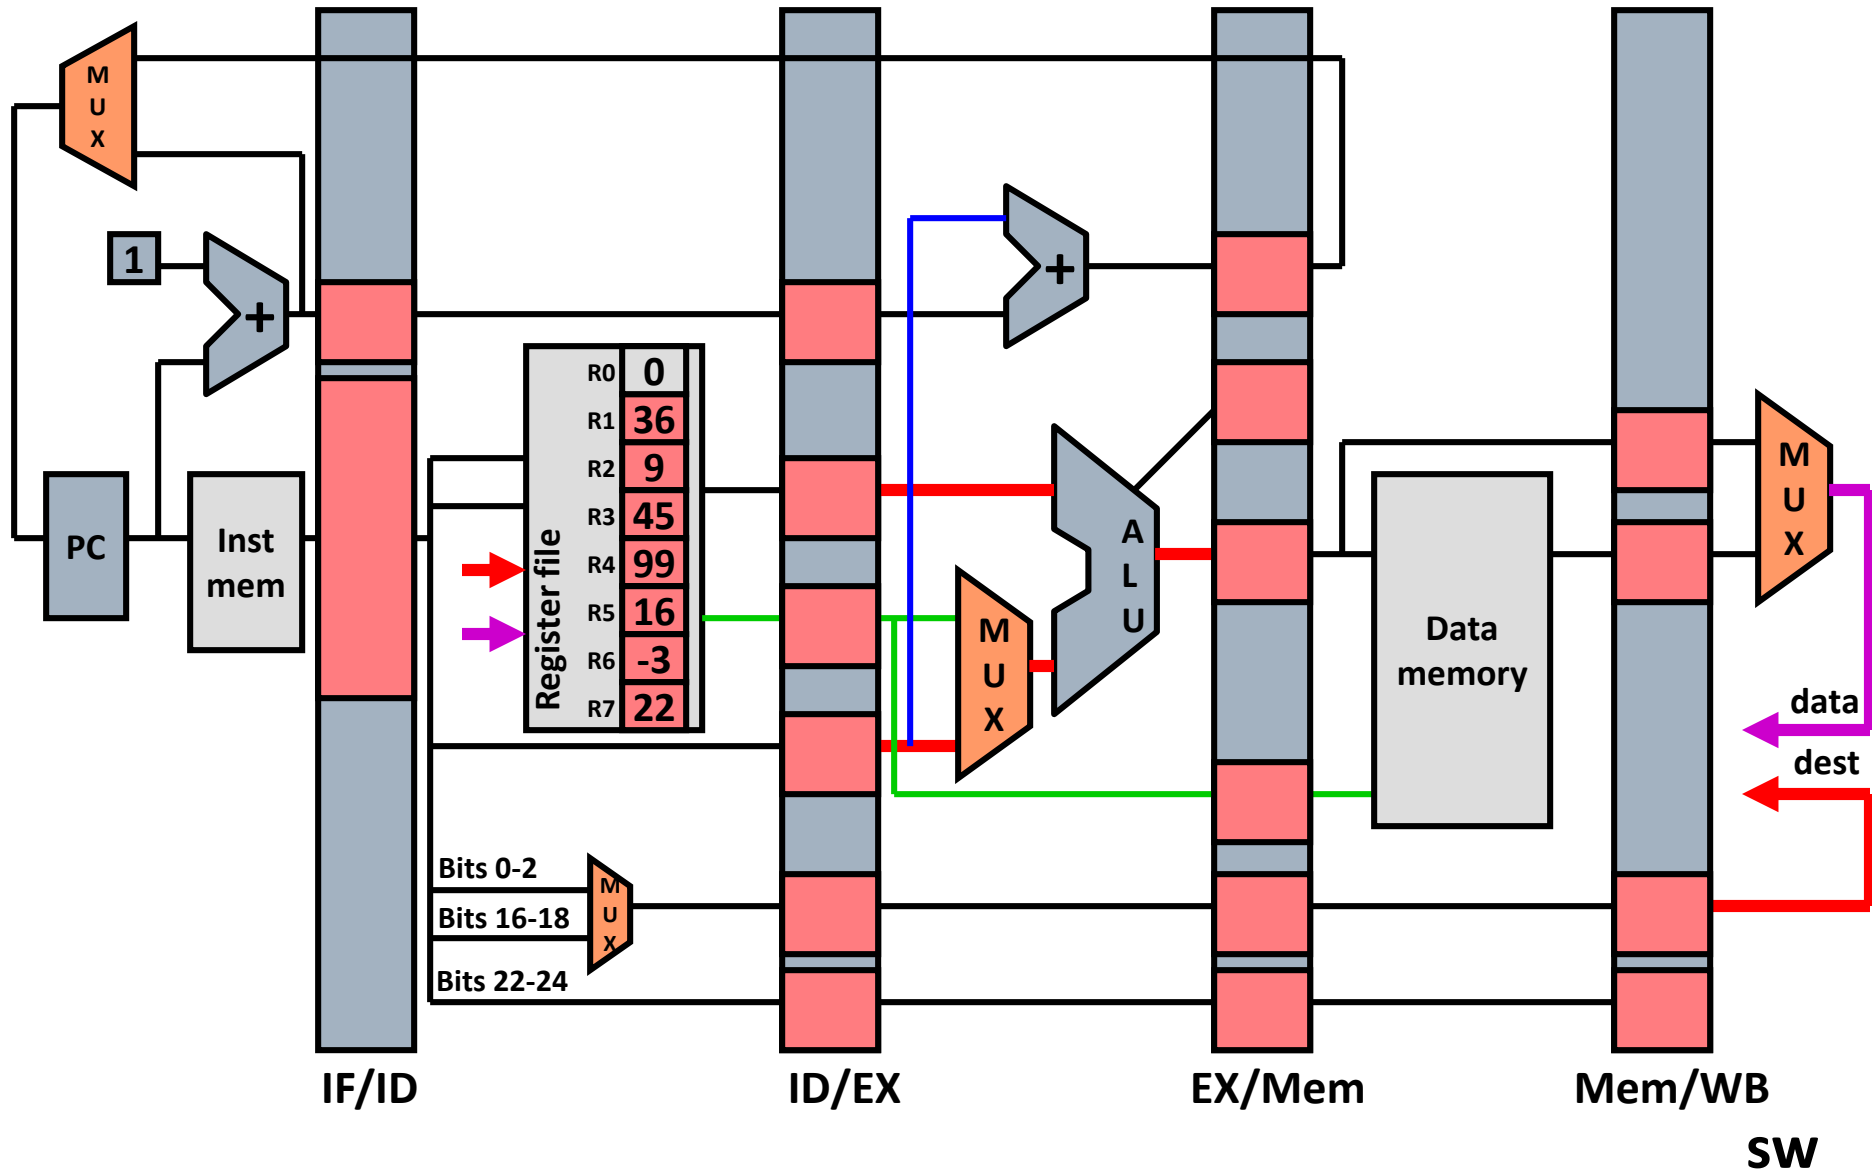
Time 7 – no more instructions



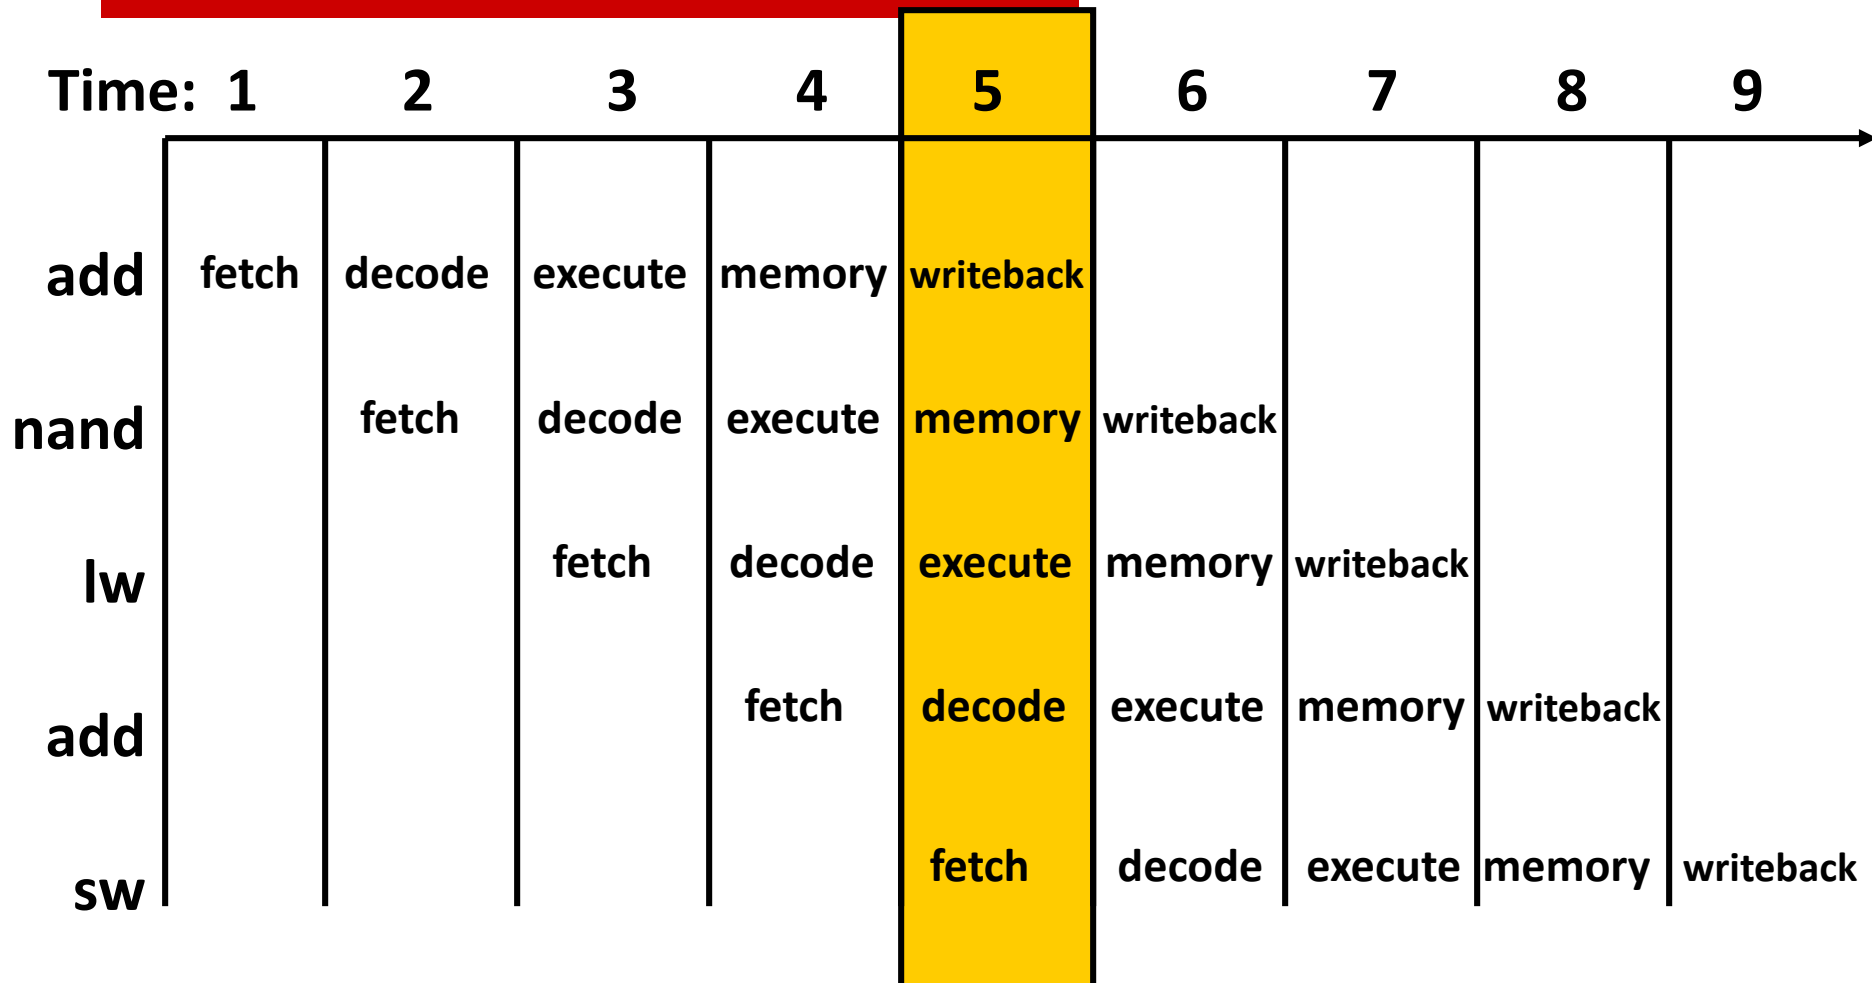
Time 8 – no more instructions



Time 9 – no more instructions



Time graphs (a.k.a. pipe trace)



A vertical slice reports the entire activity of the pipeline at time 5

What can go wrong?

- ❑ **Data hazards:** since register reads occur in stage 2 and register writes occur in stage 5 it is possible to read the wrong value if it is about to be written.
- ❑ **Control hazards:** A branch instruction may change the PC, but not until stage 4. What do we fetch before that?
- ❑ **Exceptions:** How do you handle exceptions in a pipelined processor with 5 instructions in flight?

- ❑ Next Lecture: Data Hazards